

**THE ONE ABOUT
KAFKA
LILY MARA**



OneSignal

Refactoring to Rust

Lily Mara
Joel Holmes

MEAP

MANNING



ONCE UPON A TIME

ONCE UPON A TIME

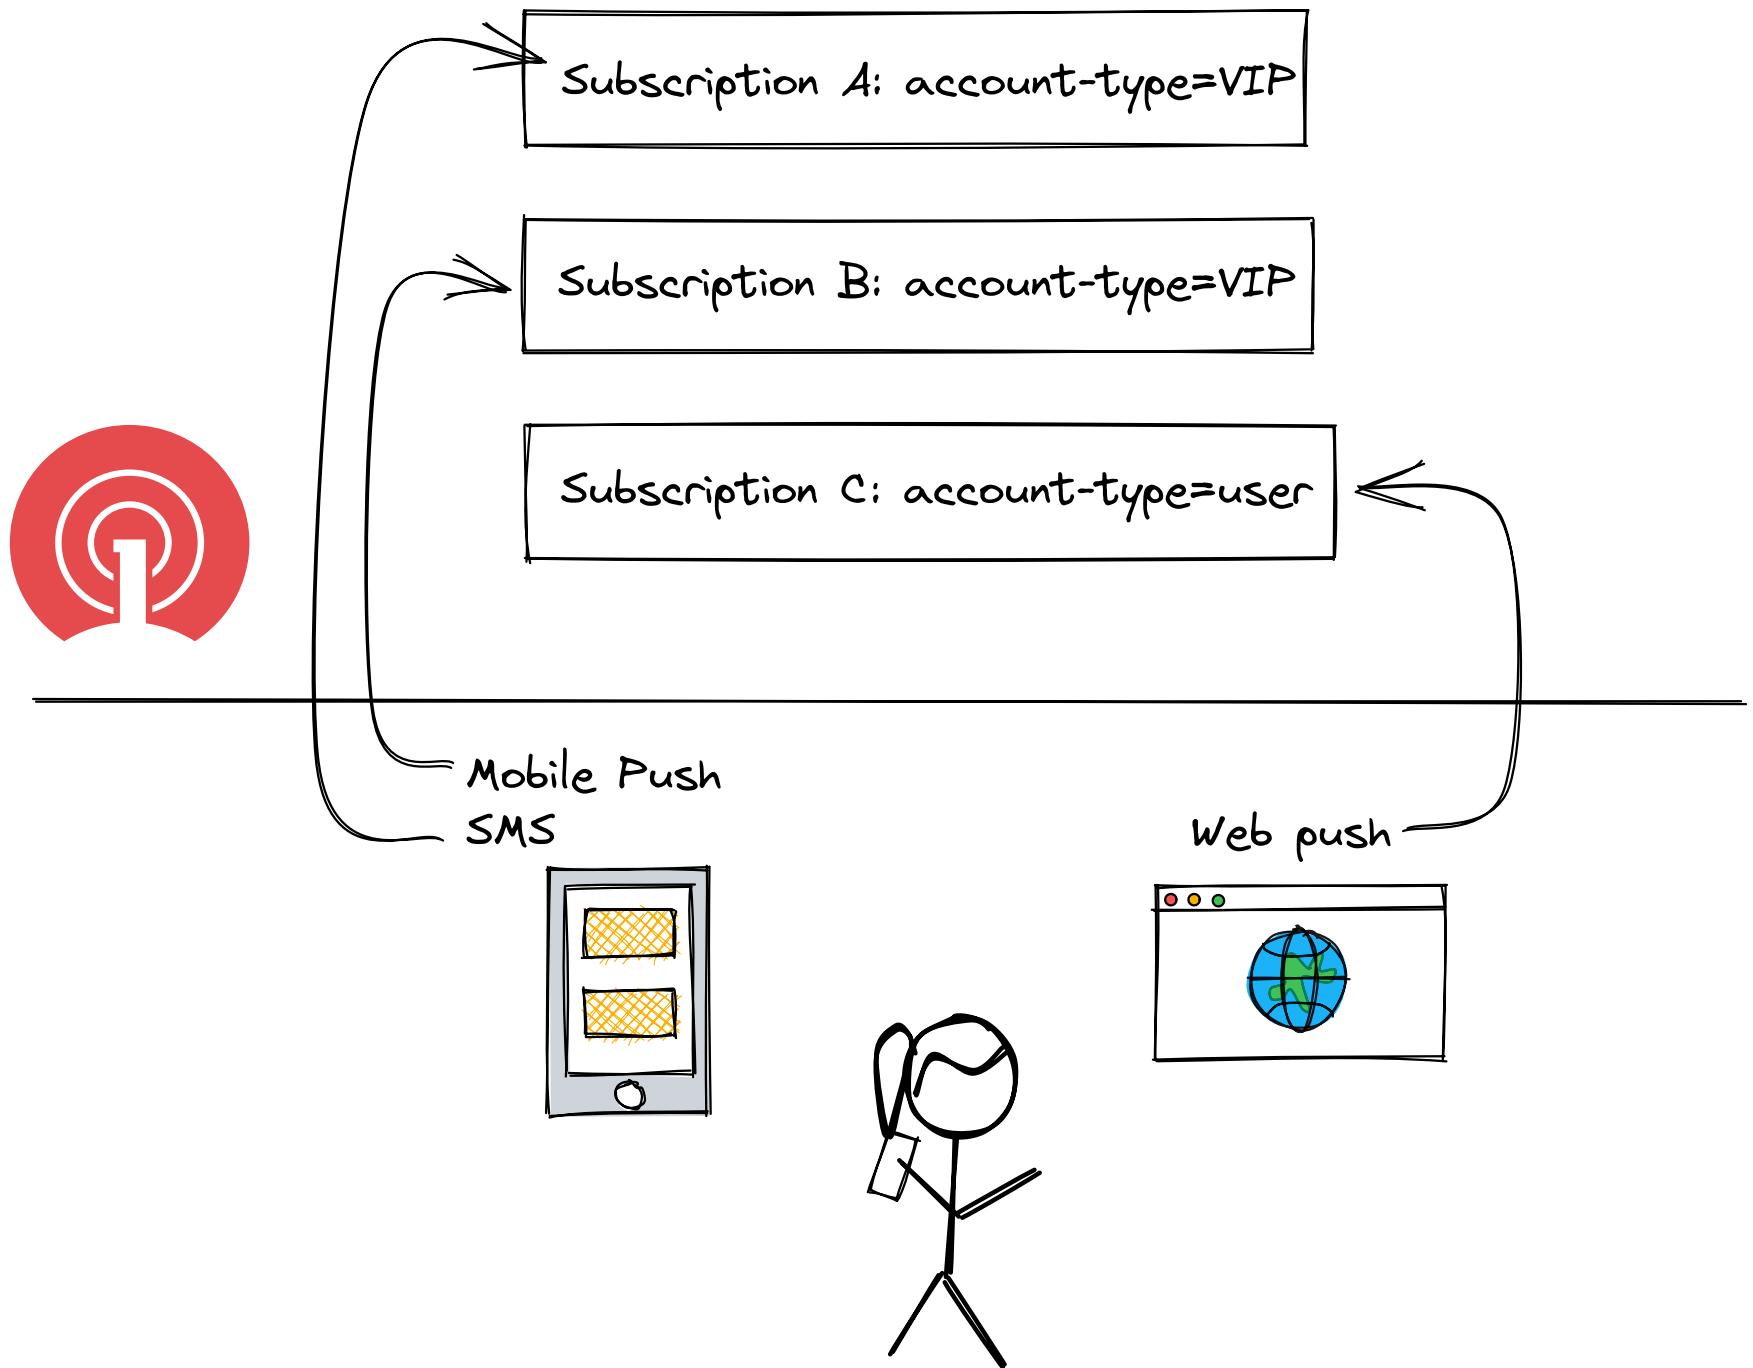
- 8B notifications/day

ONCE UPON A TIME

- 8B notifications/day
- 10 backend engineers

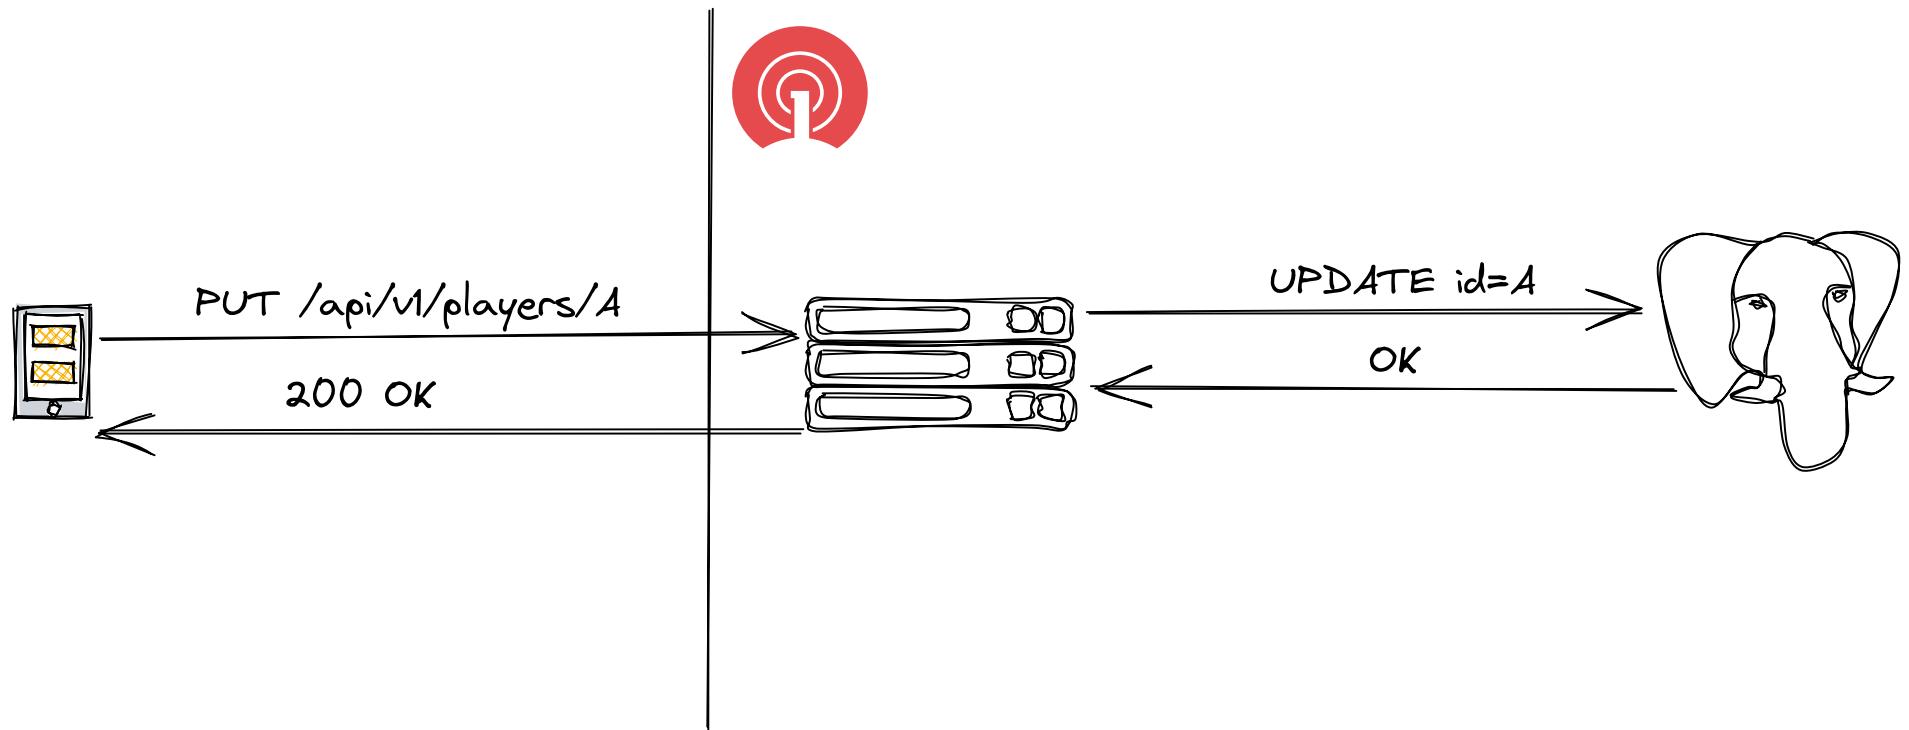
ONCE UPON A TIME

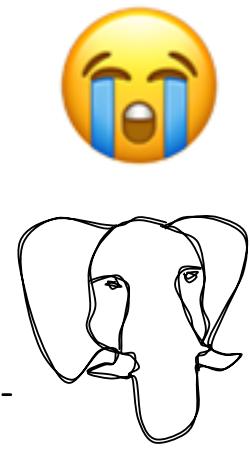
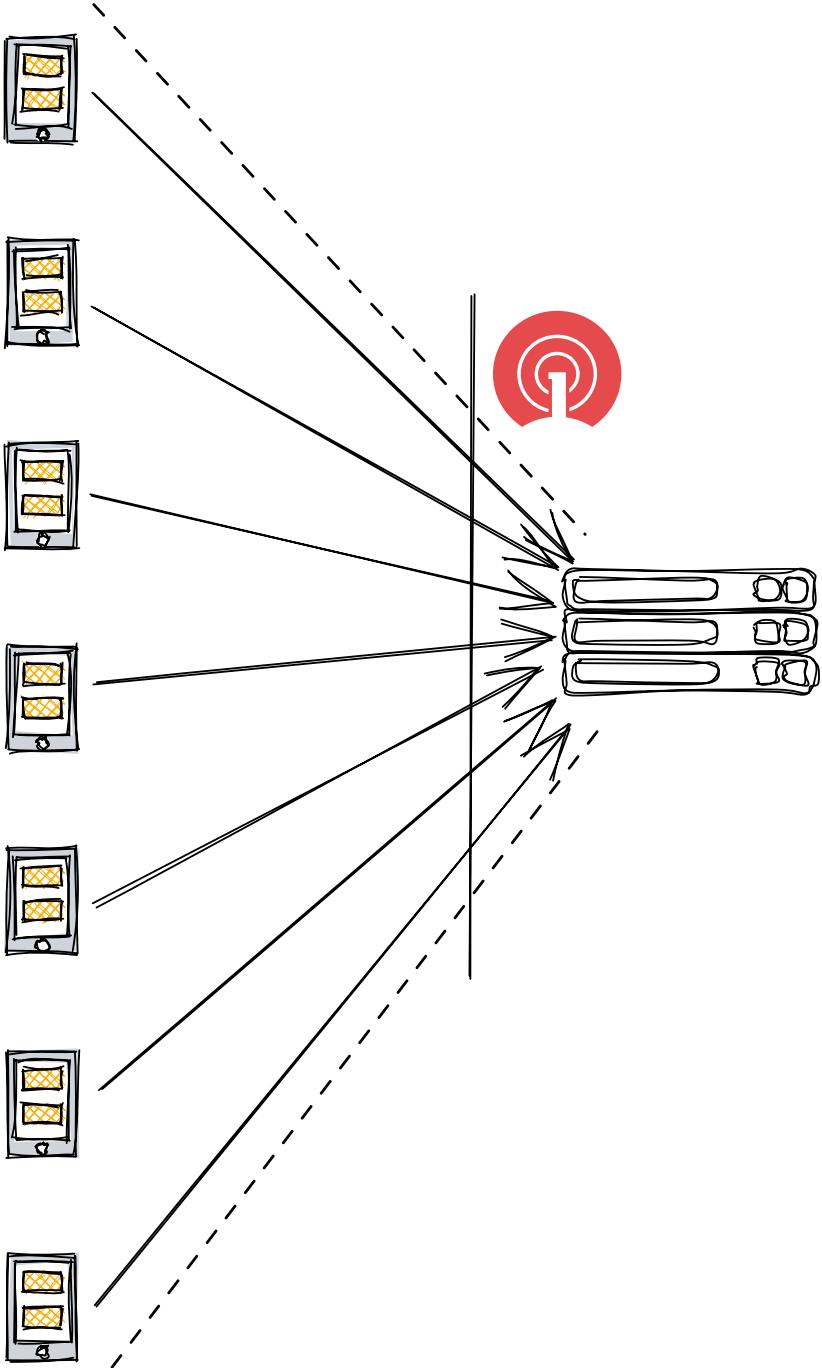
- 8B notifications/day
- 10 backend engineers
- Make simplifying assumptions



API

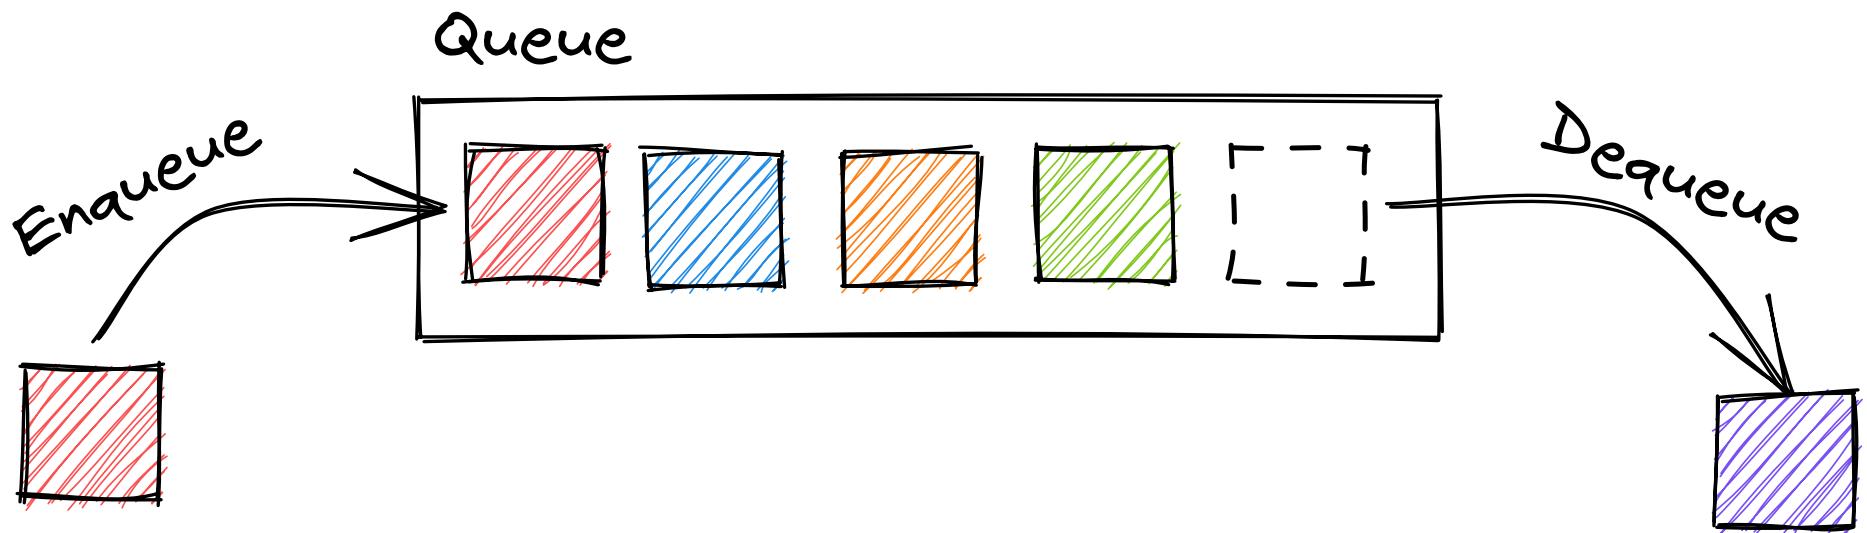
```
PUT https://onesignal.com/api/v1/players/{SUBSCRIPTION_ID}
{
    "app_id": "{APP_ID}",
    "tags": {
        "first_name": "Jon",
        "last_name": "Smith",
    }
}
```



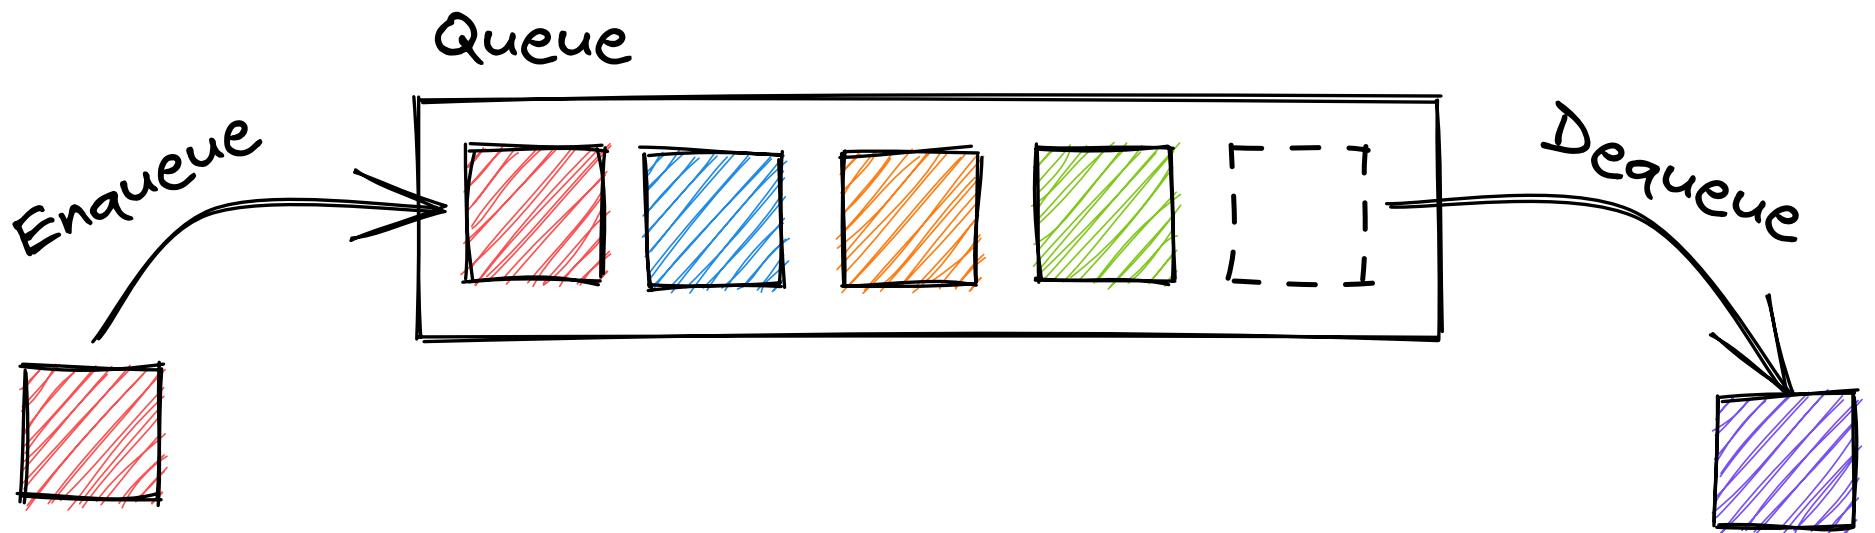


QUEUE

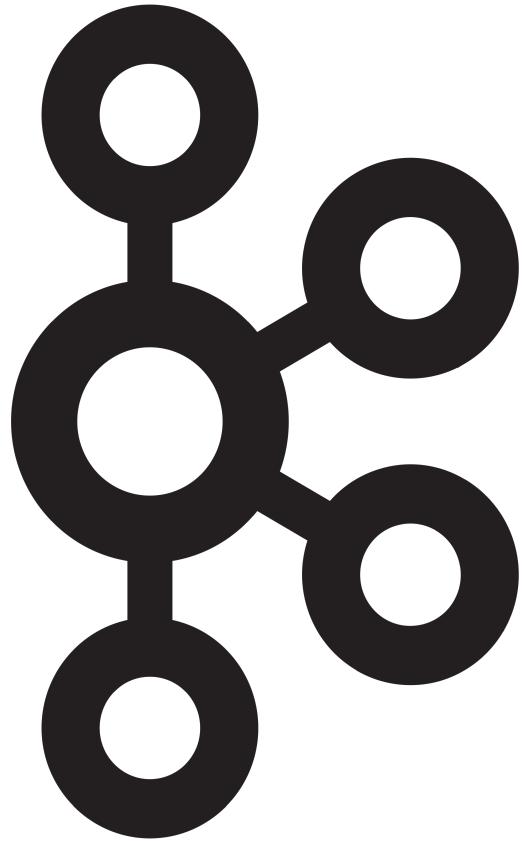
QUEUE



QUEUE

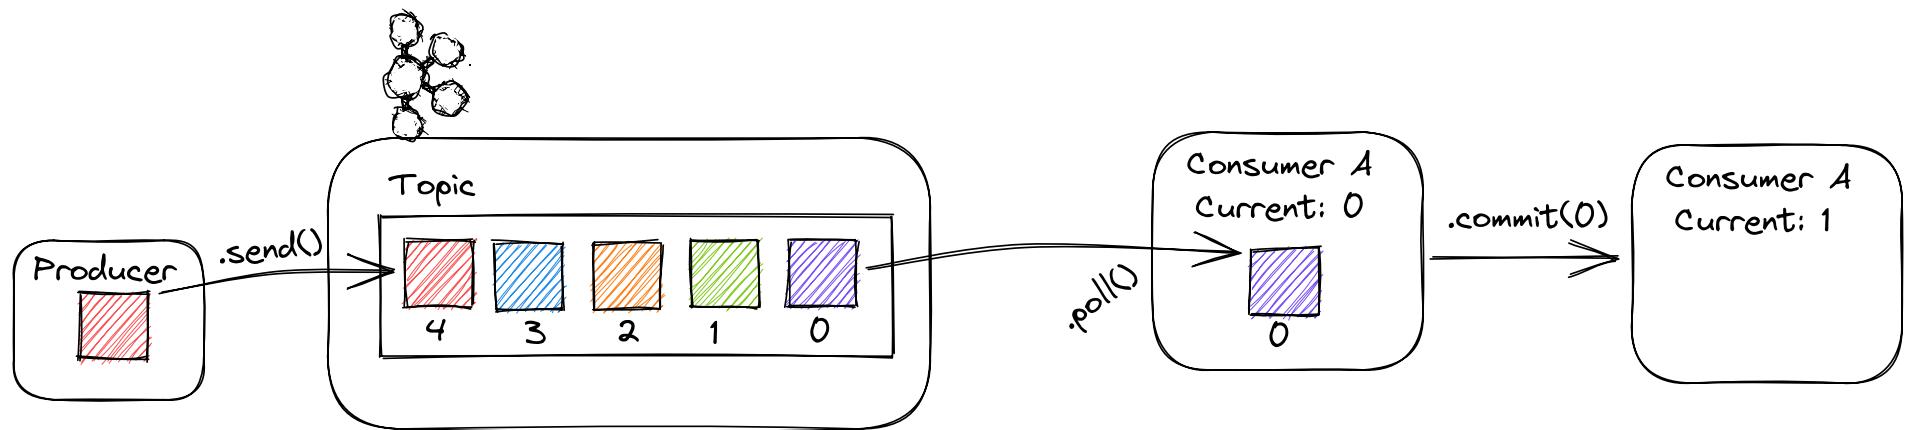


Metric: Lag=4

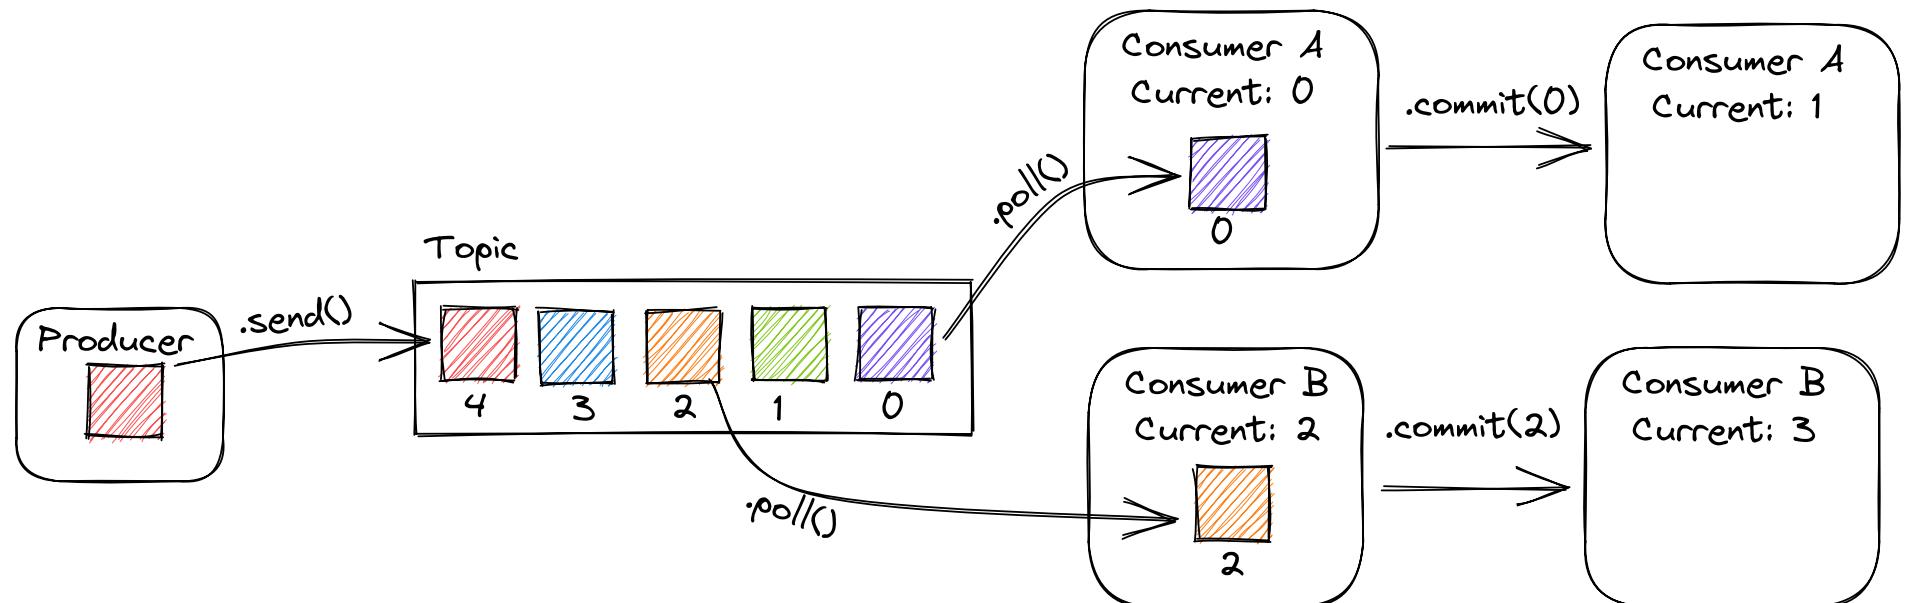


kafka

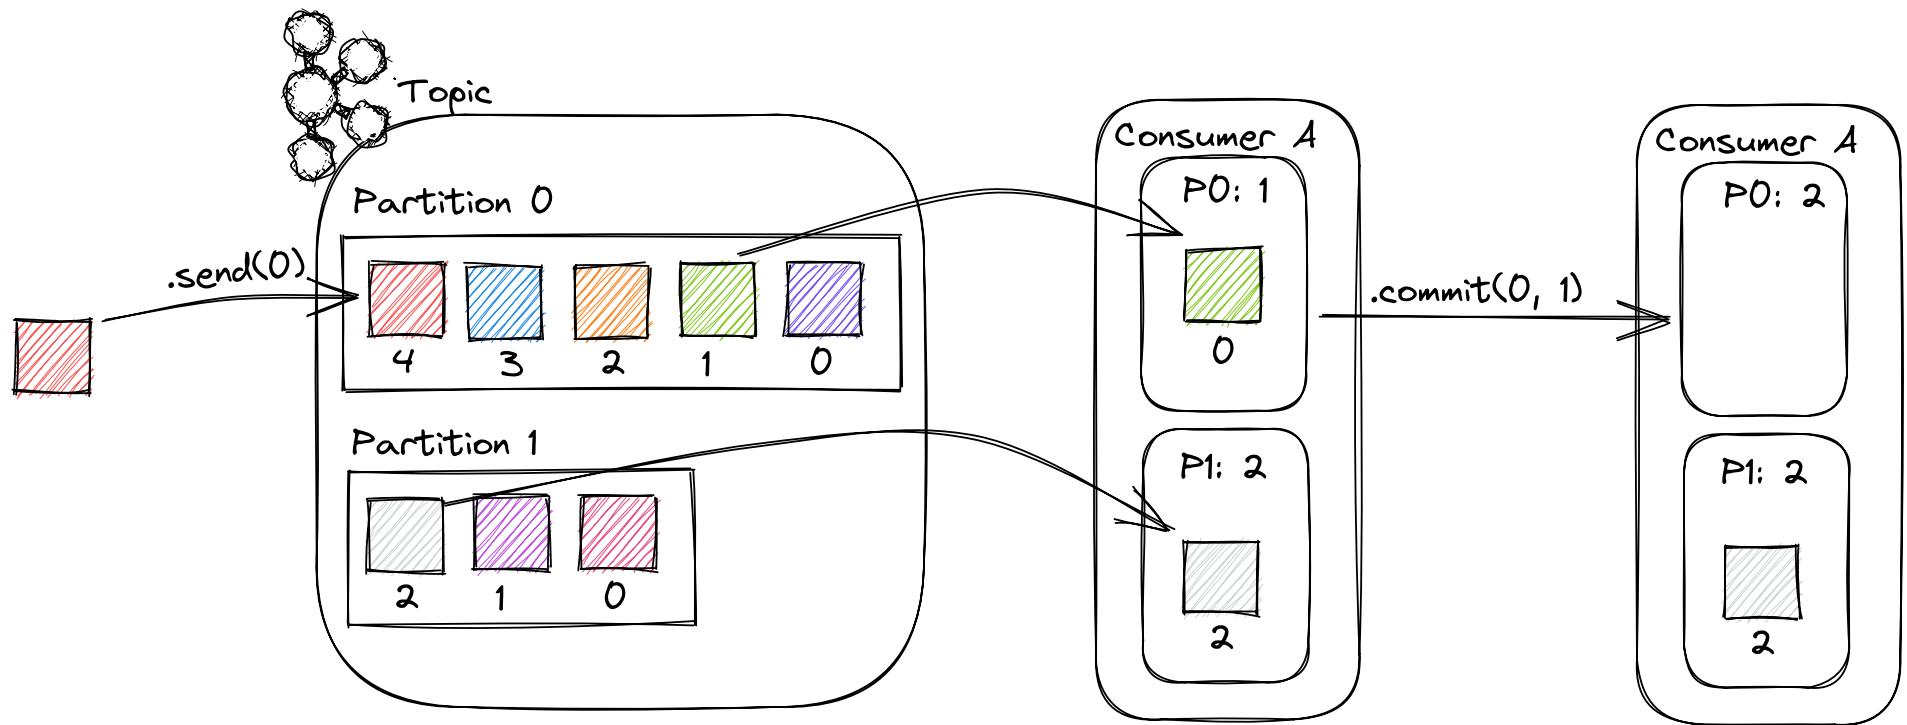
KAFKA PIPELINE



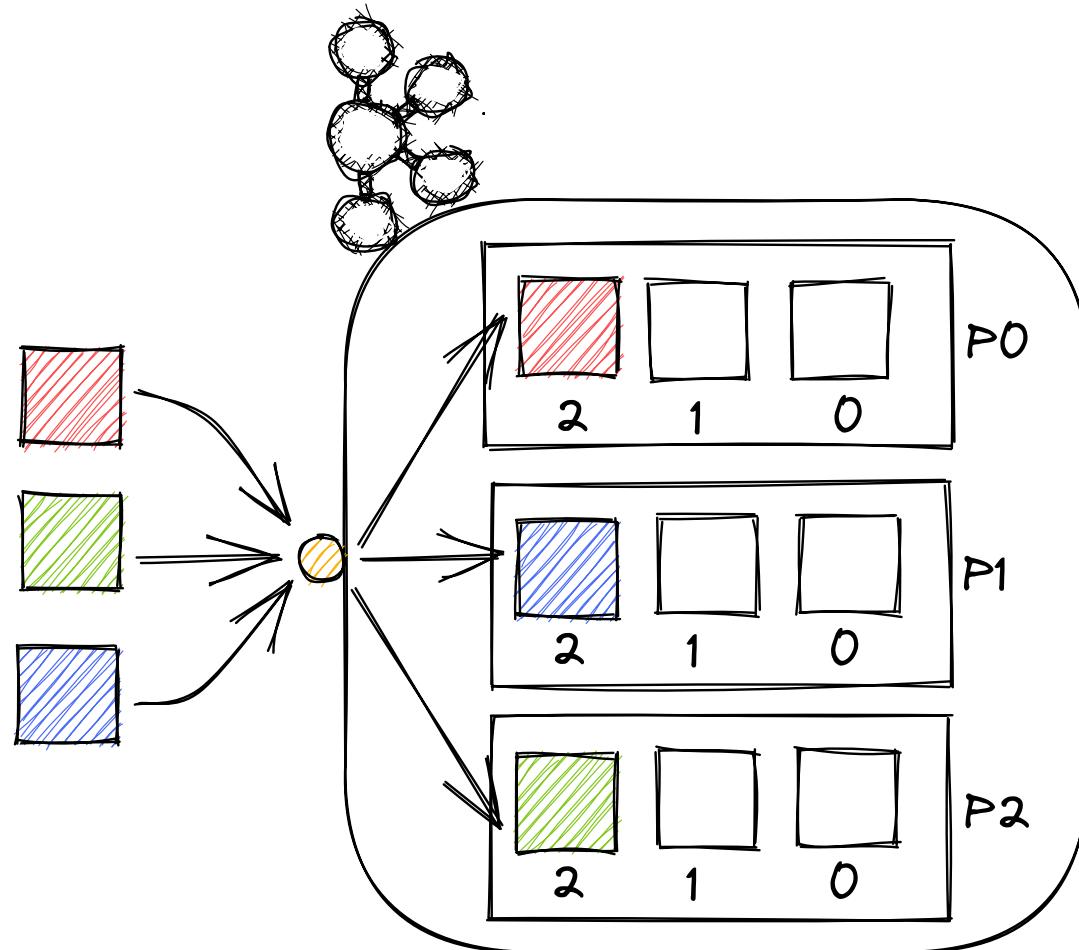
CONSUMER



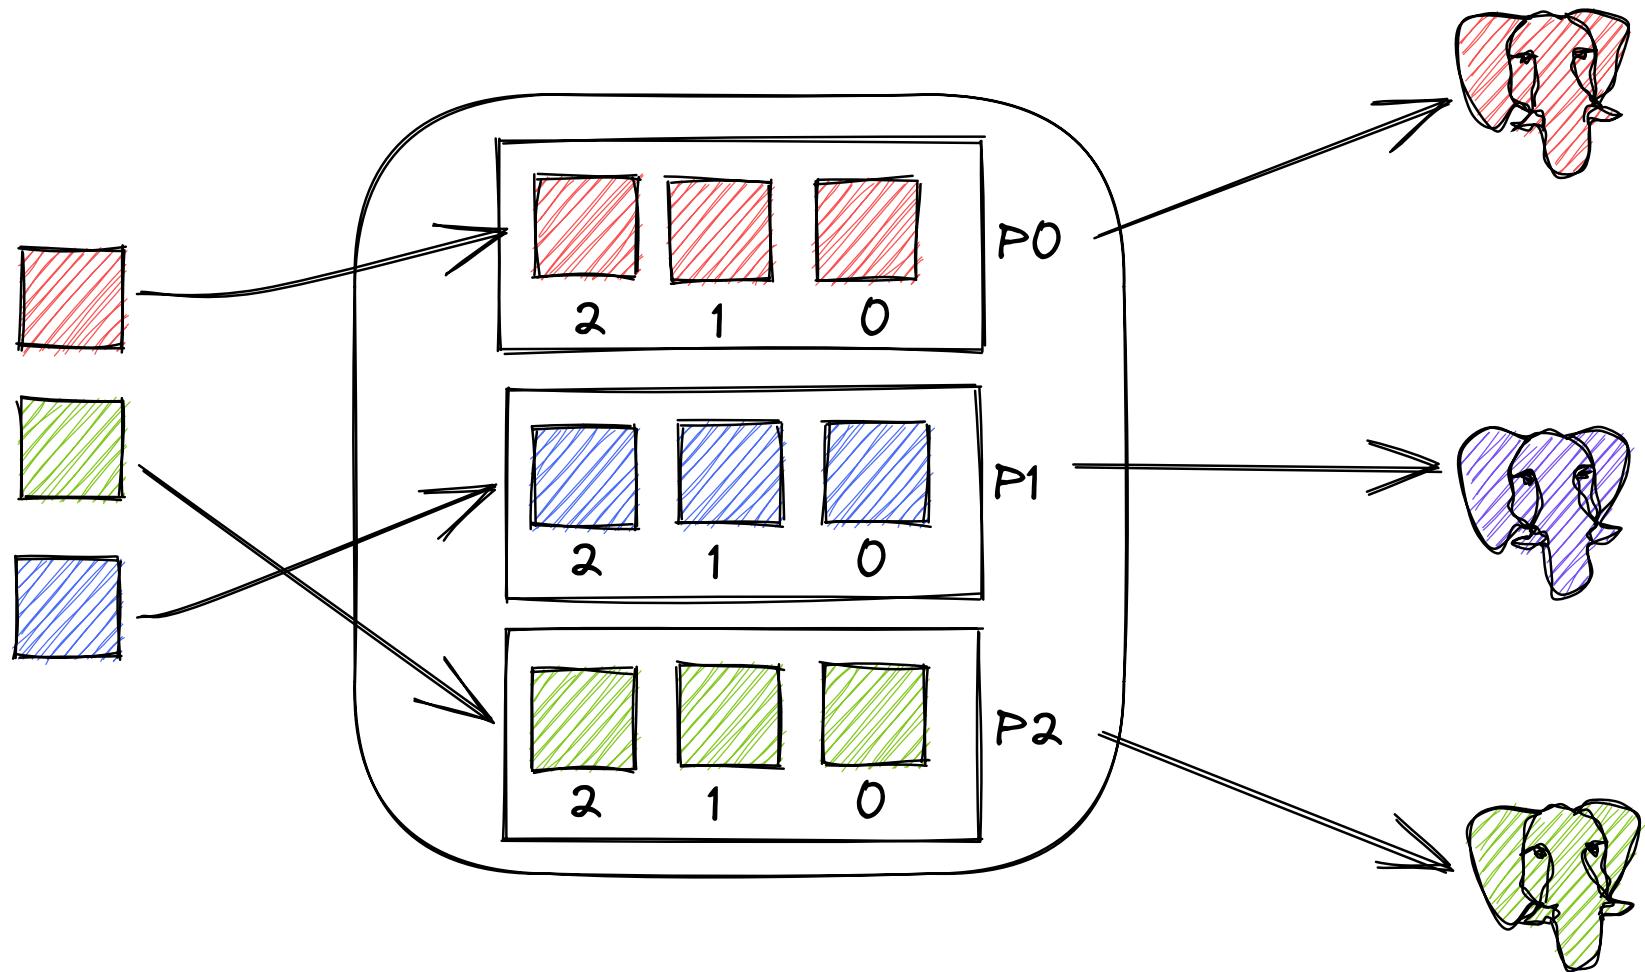
PARTITION



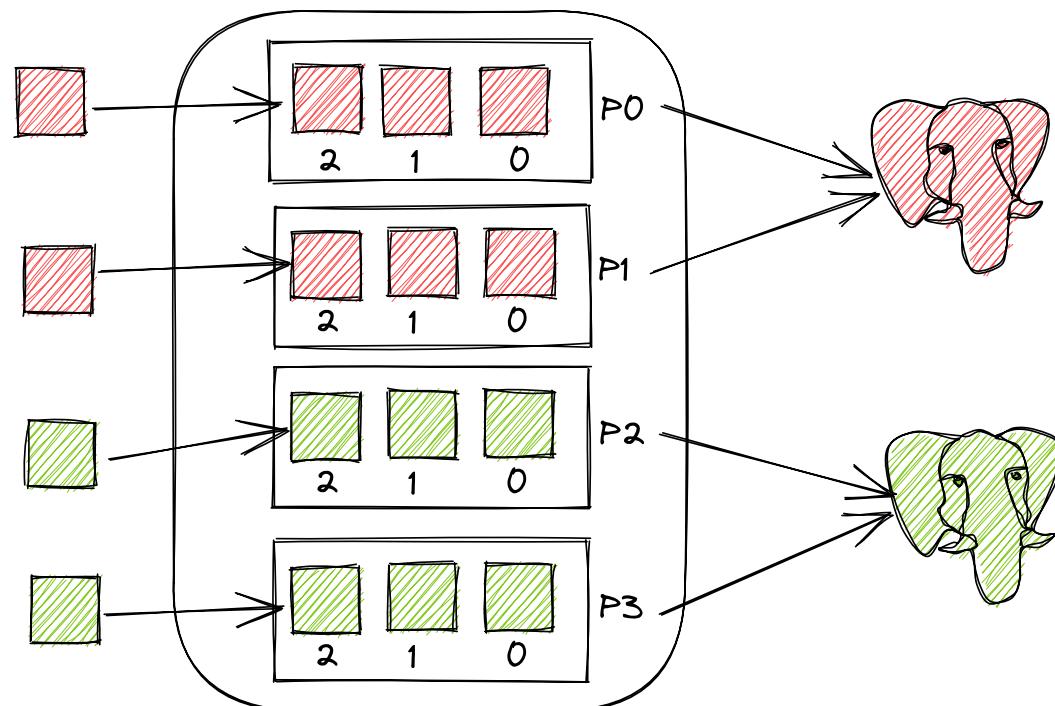
ROUND-ROBIN



EXPLICIT



CONCURRENT WRITES



ISSUES

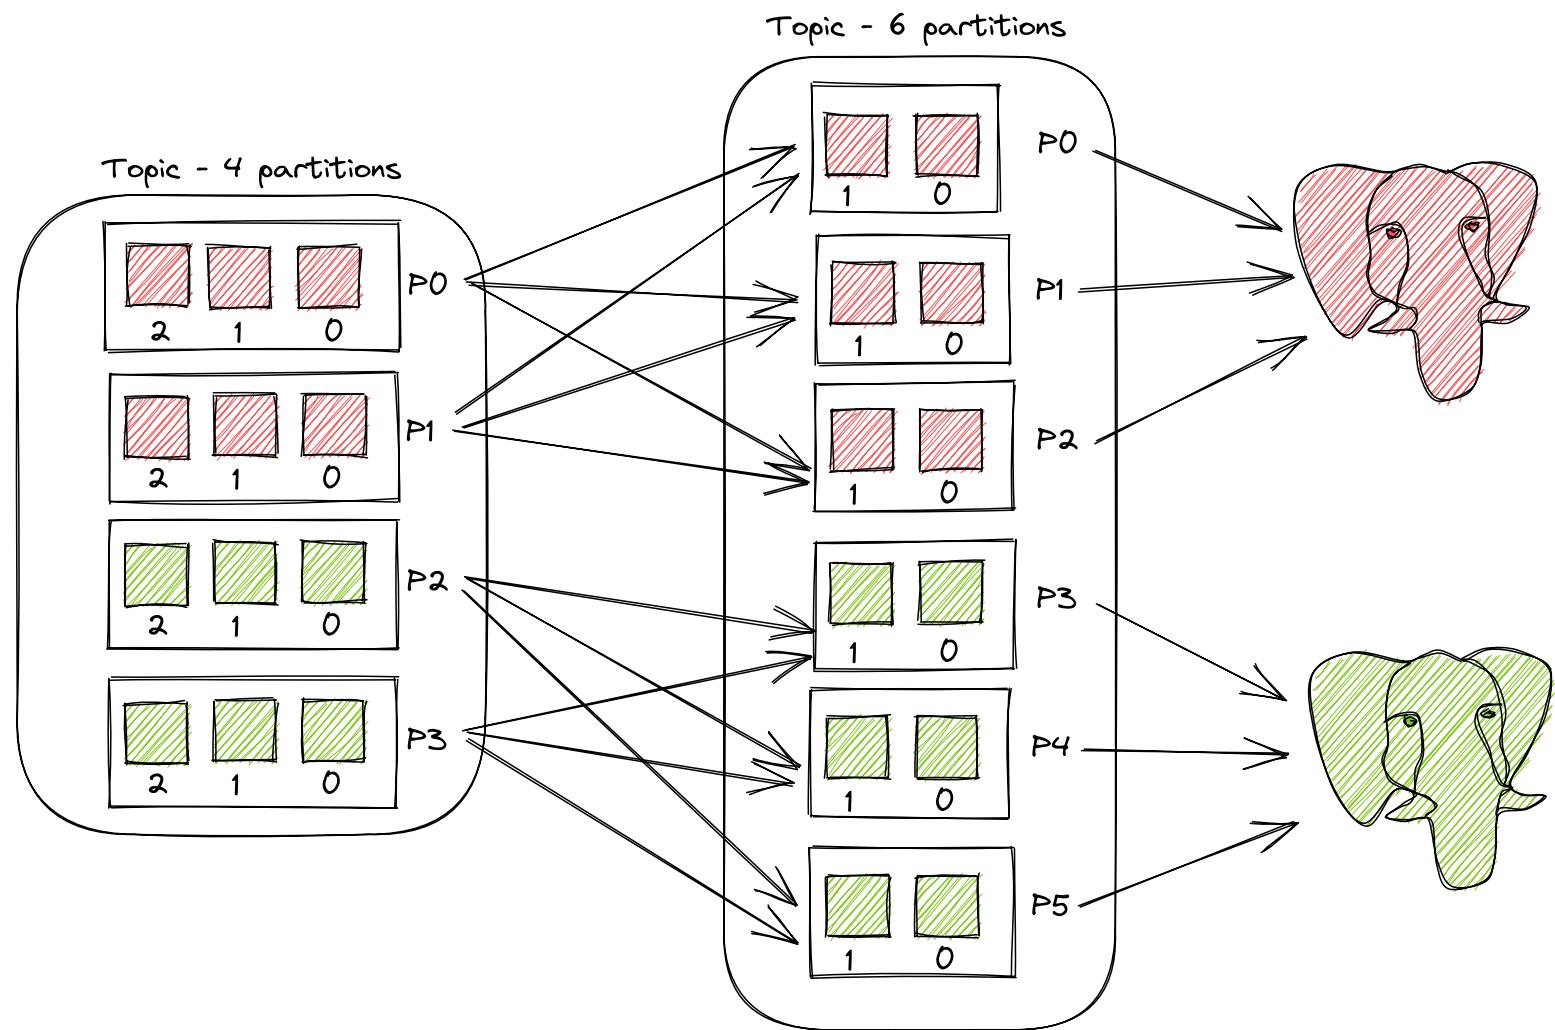
ISSUES

- Inflexible

ISSUES

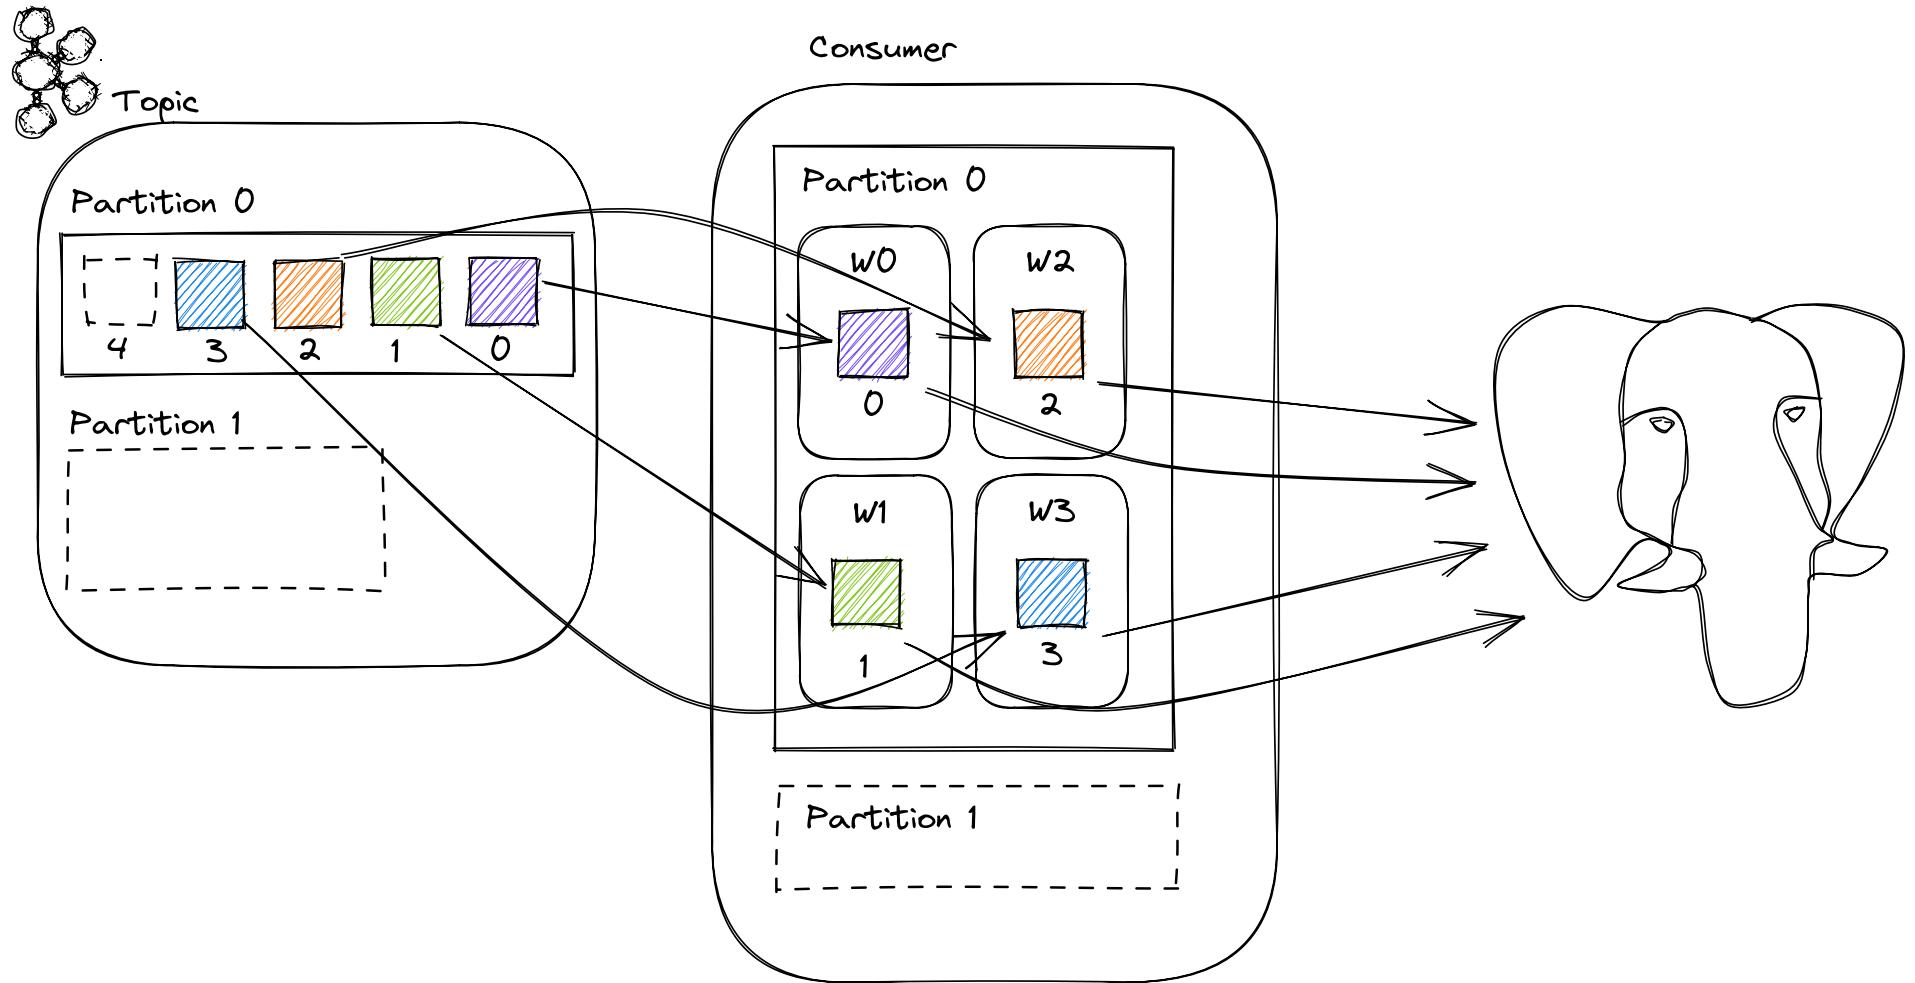
- Inflexible
- Kafka repartitioning

REPARTITIONING

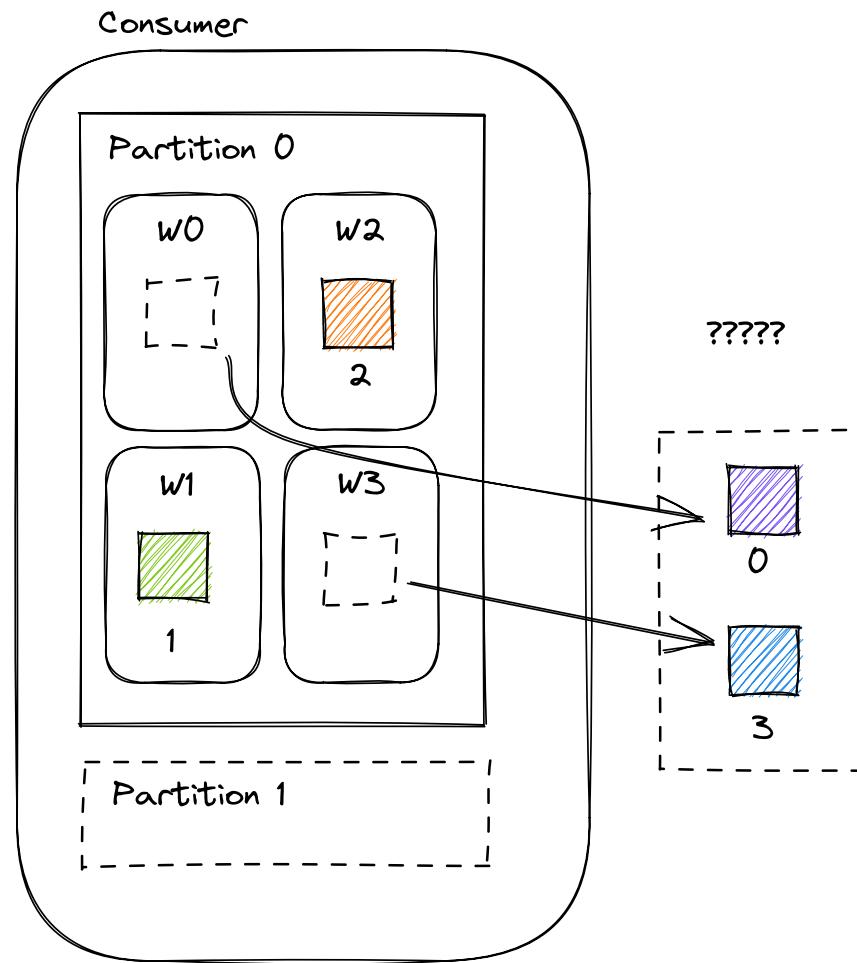


SUBPARTITION PROCESSING

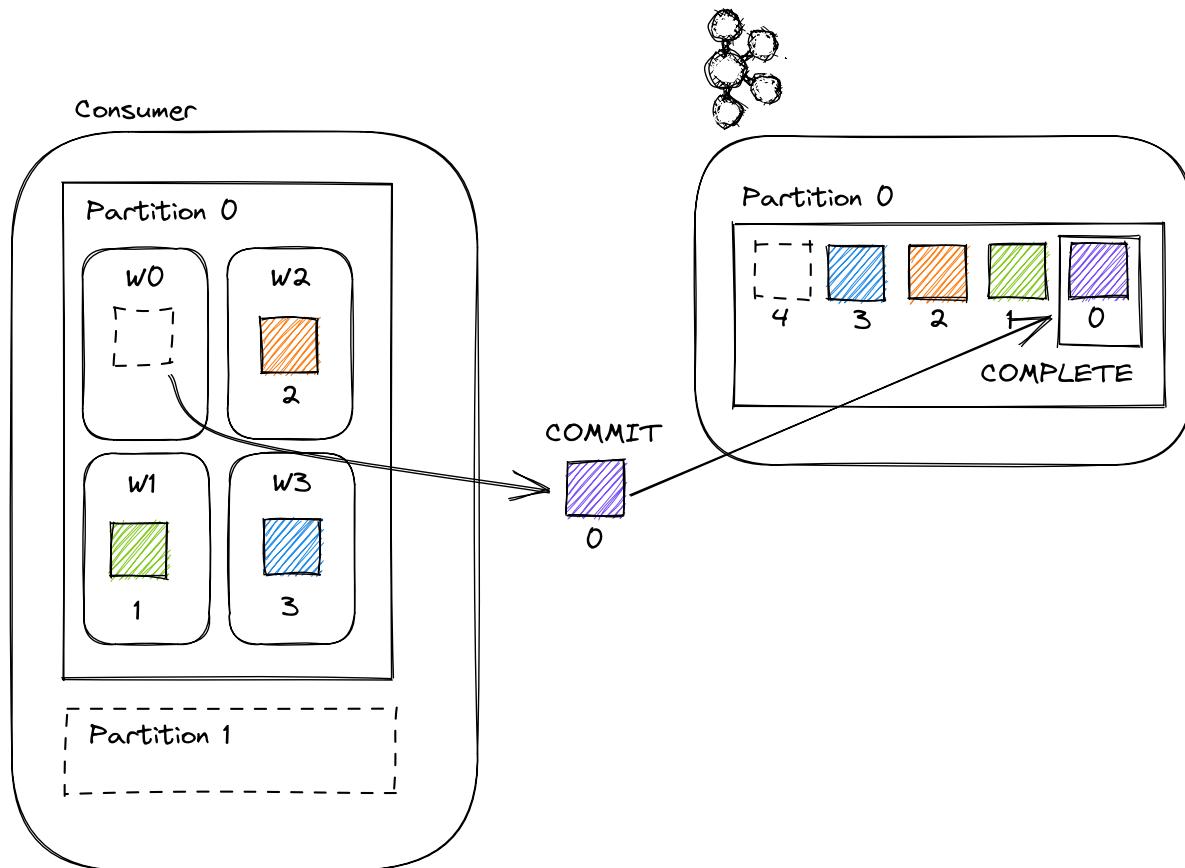
SUBPARTITION PROCESSING



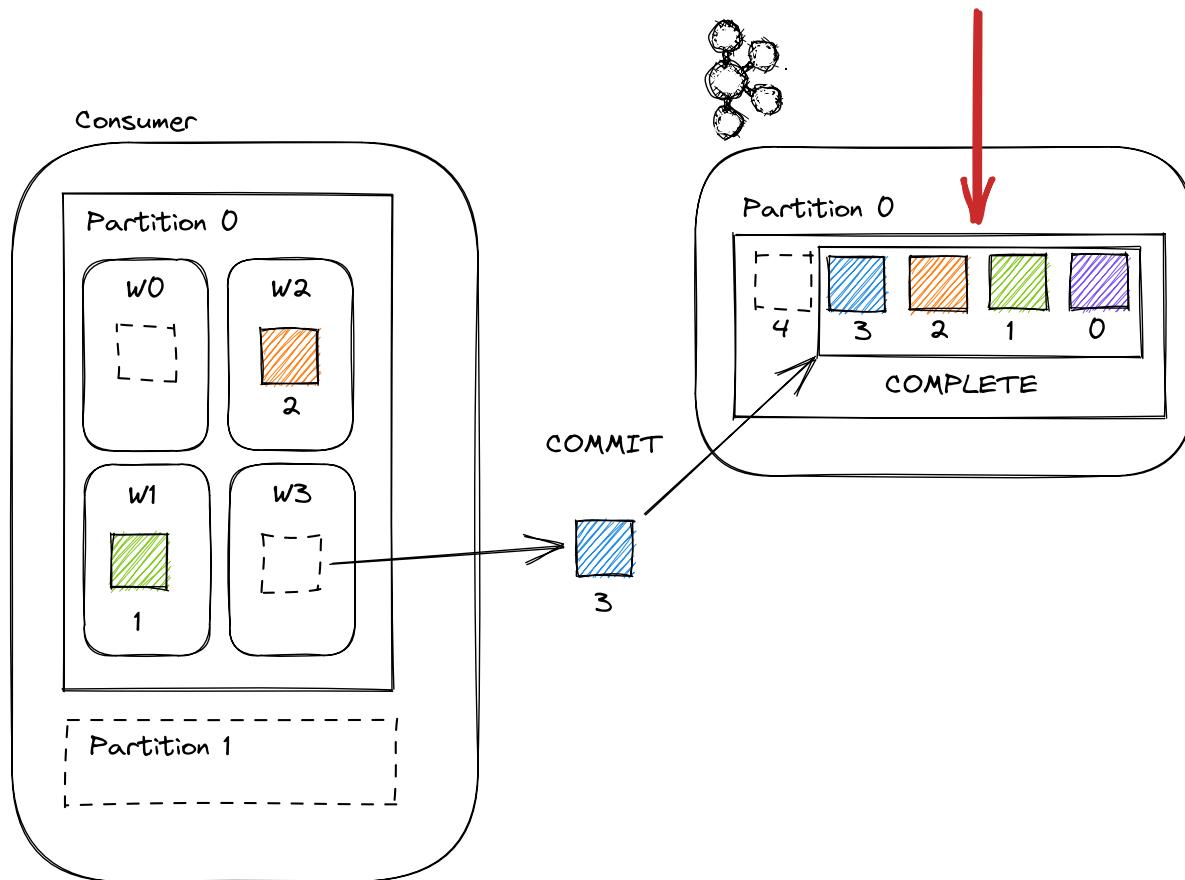
ISSUE



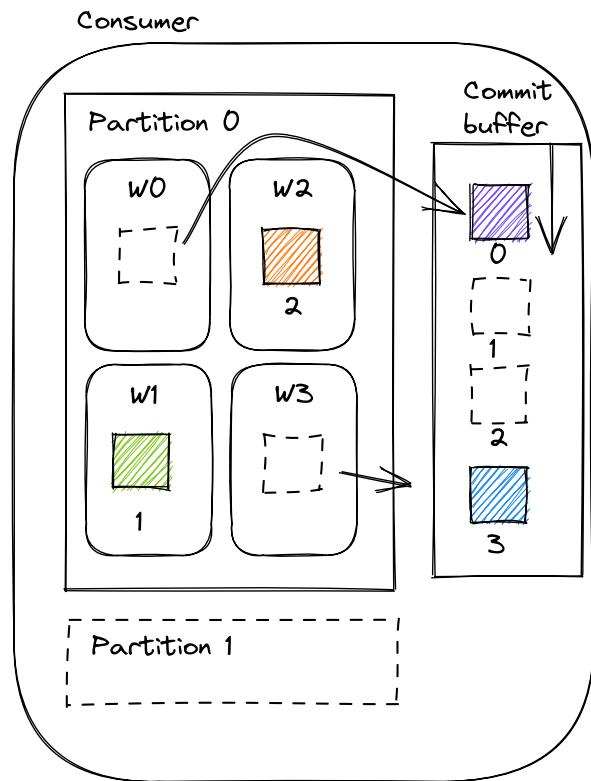
COMMIT(0)



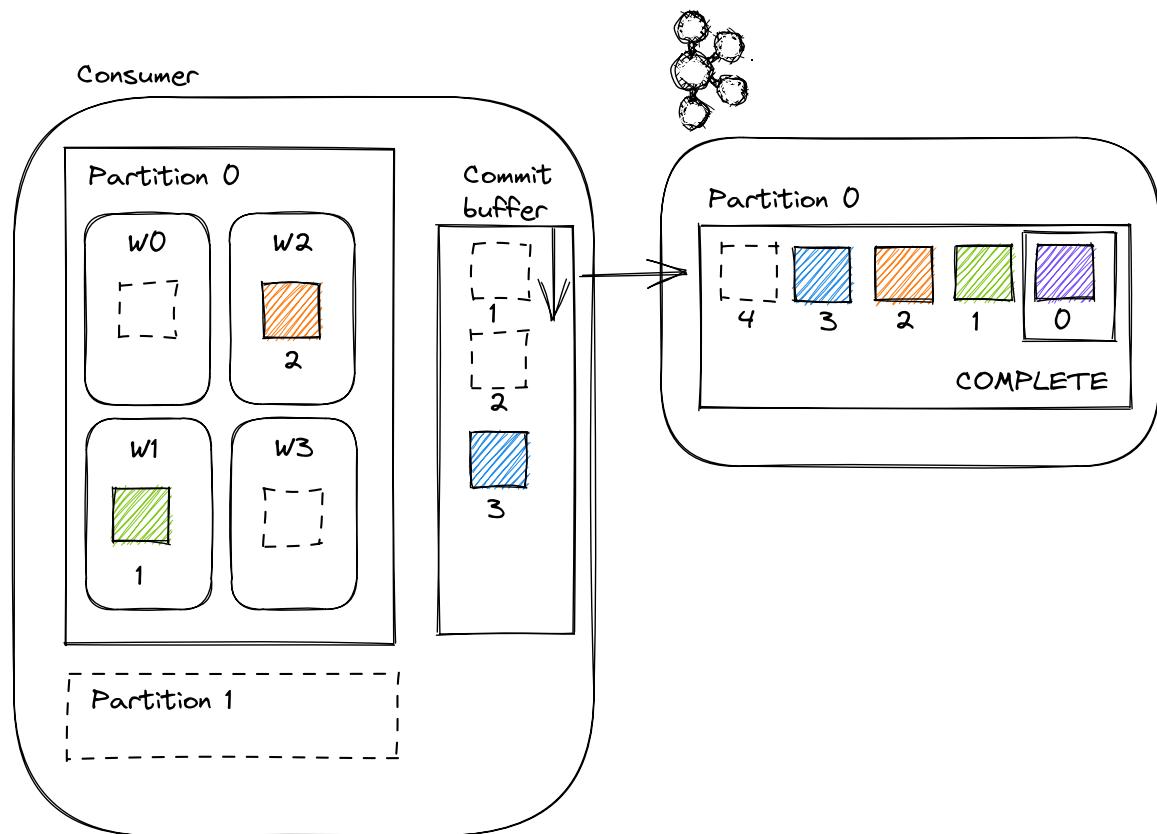
COMMIT(3)



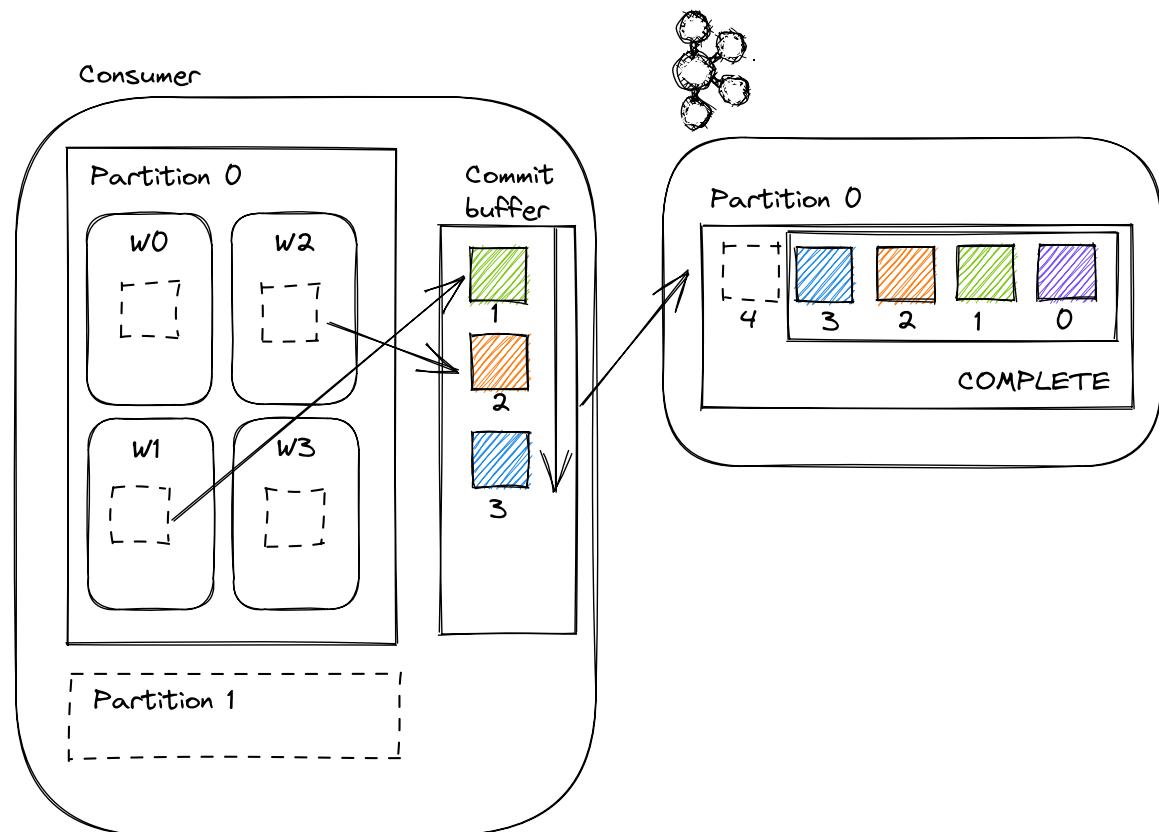
SOLUTION



COMMIT(0)



COMMIT(3)



CONCESSION

CONCESSION

- at-least-once delivery

CONCESSION

- at-least-once delivery
- messages will be replayed

CONCESSION

- at-least-once delivery
- messages will be replayed
- design around this

REVIEW

REVIEW

- kafka topic

REVIEW

- kafka topic
- contains partitions - queues

REVIEW

- kafka topic
- contains partitions - queues
- message has incrementing offset

REVIEW

- kafka topic
- contains partitions - queues
- message has incrementing offset
- producers enqueue

REVIEW

- kafka topic
- contains partitions - queues
- message has incrementing offset
- producers enqueue
- consumers dequeue

REVIEW

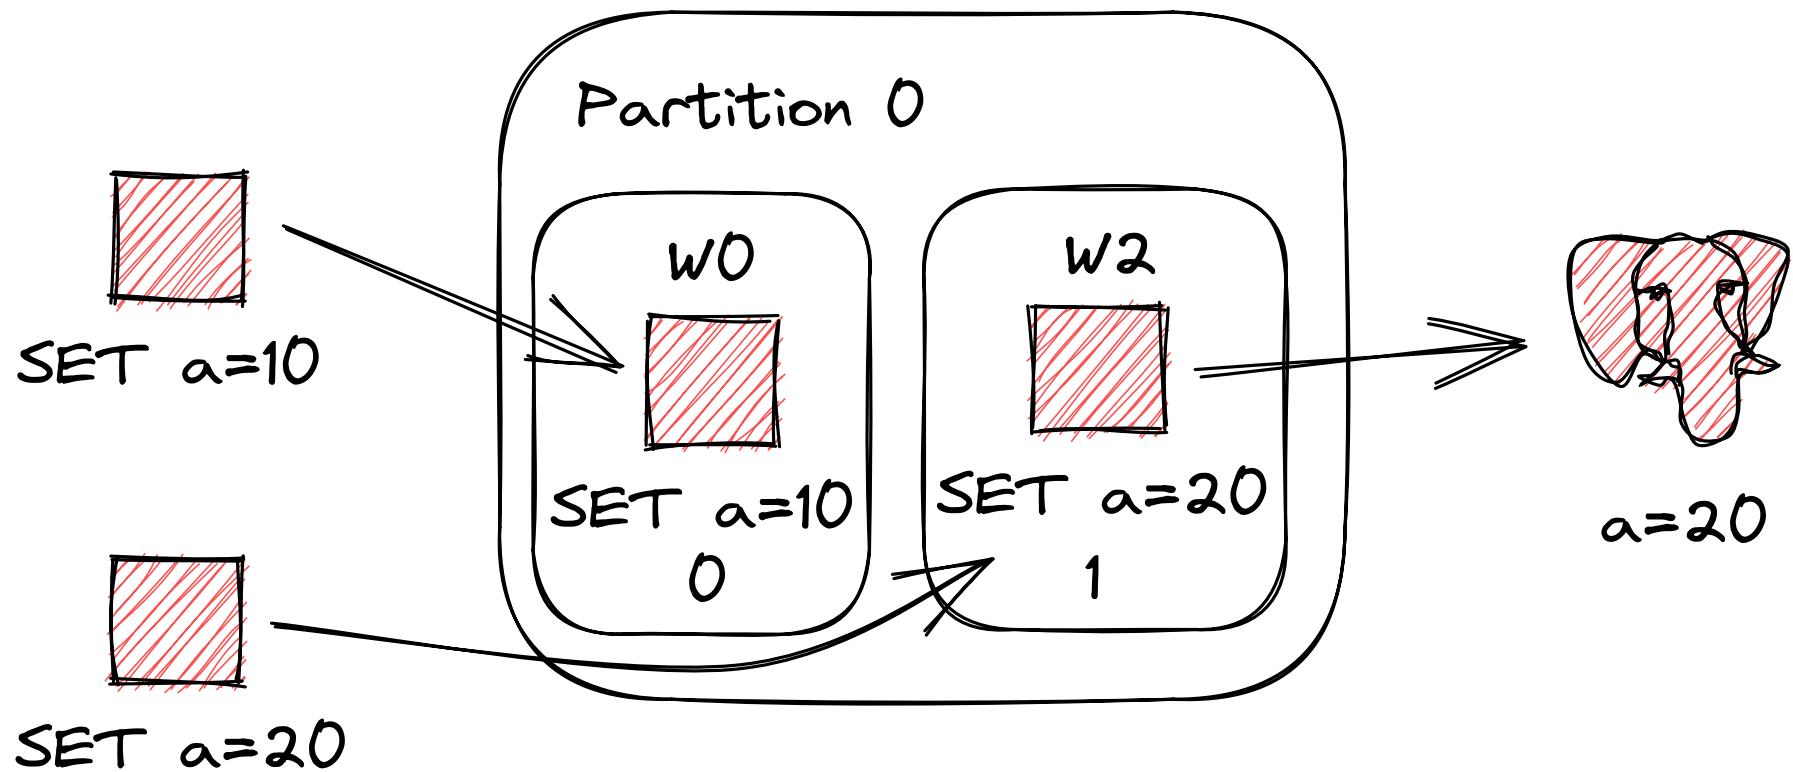
- kafka topic
- contains partitions - queues
- message has incrementing offset
- producers enqueue
- consumers dequeue
- consumer concurrency via partitioning and subpartitioning

REVIEW

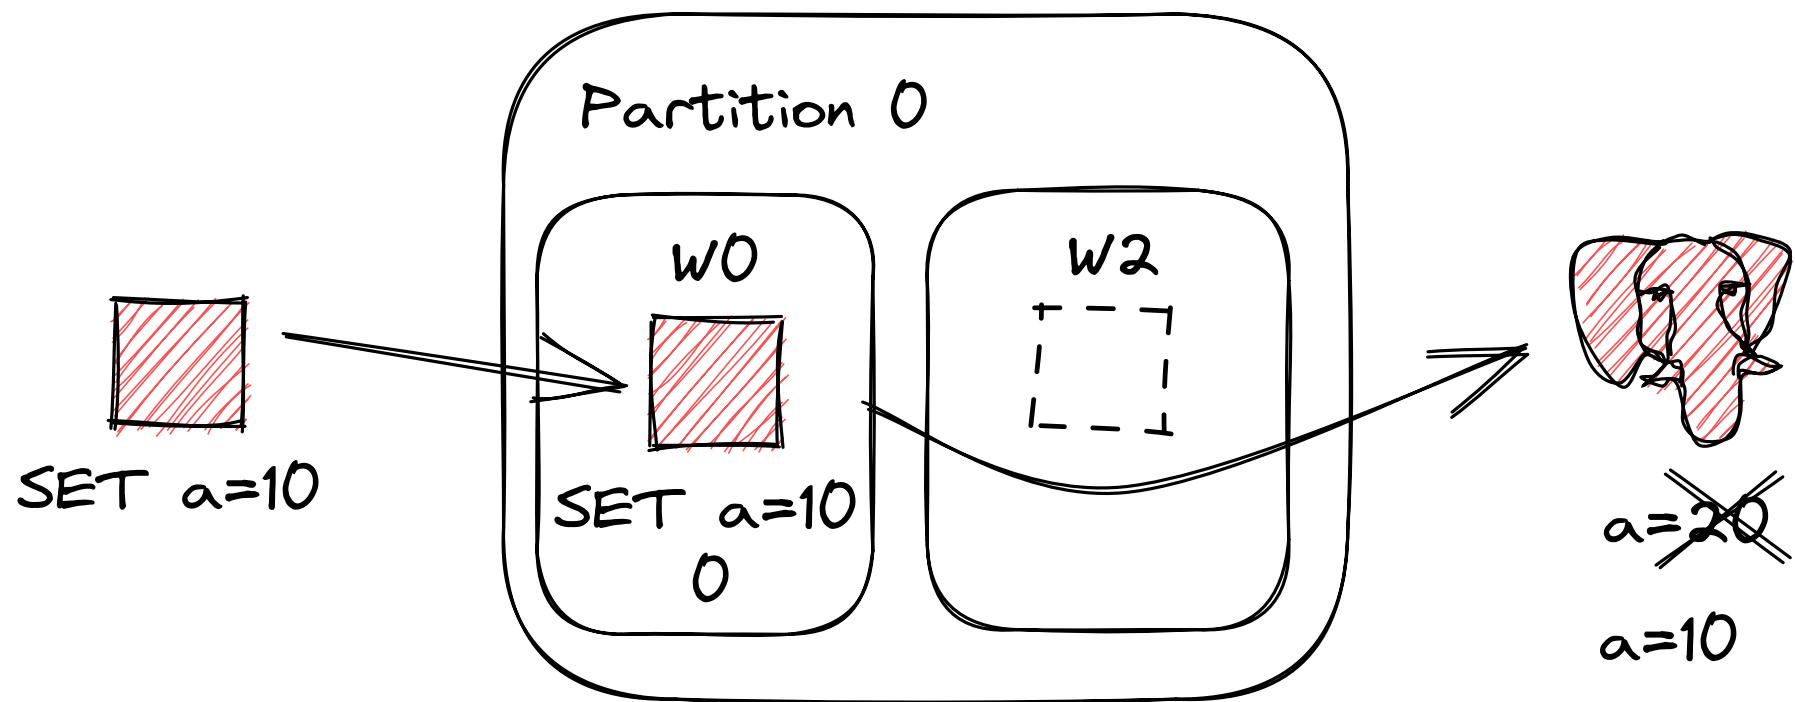
- kafka topic
- contains partitions - queues
- message has incrementing offset
- producers enqueue
- consumers dequeue
- consumer concurrency via partitioning and subpartitioning
- consumer performing PG writes

POSTGRES WRITES

CONCURRENCY



CONCURRENCY



GOALS

GOALS

- maximize concurrency

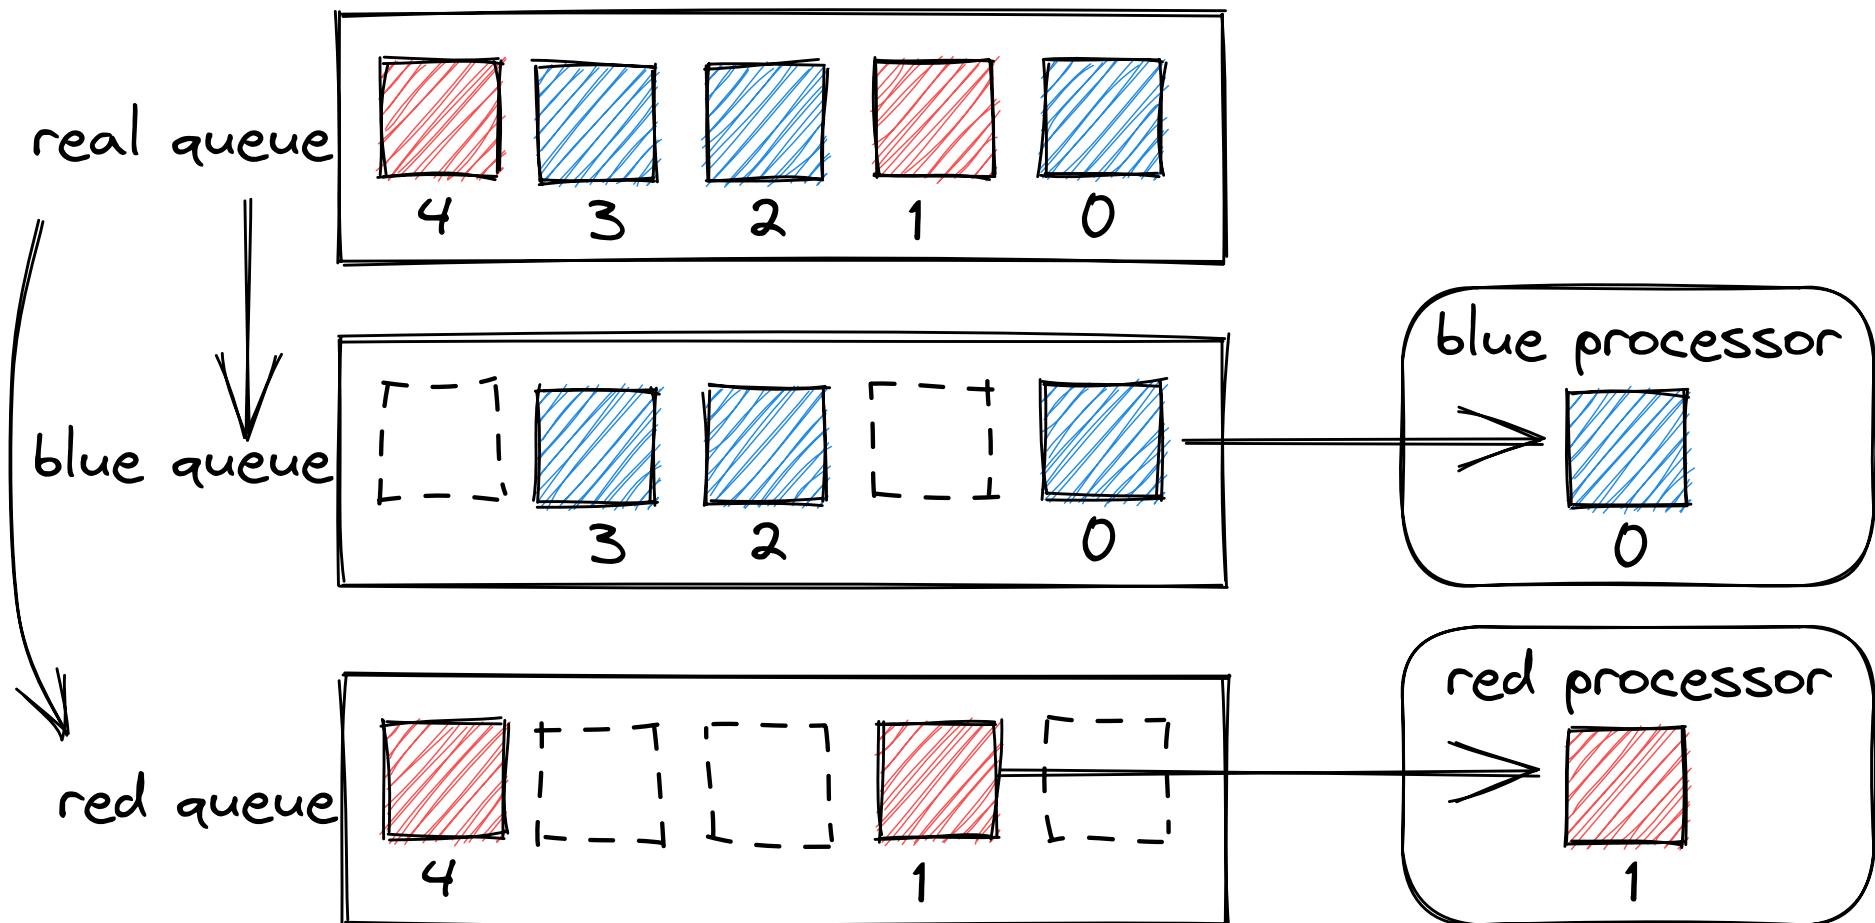
GOALS

- maximize concurrency
- minimize contention

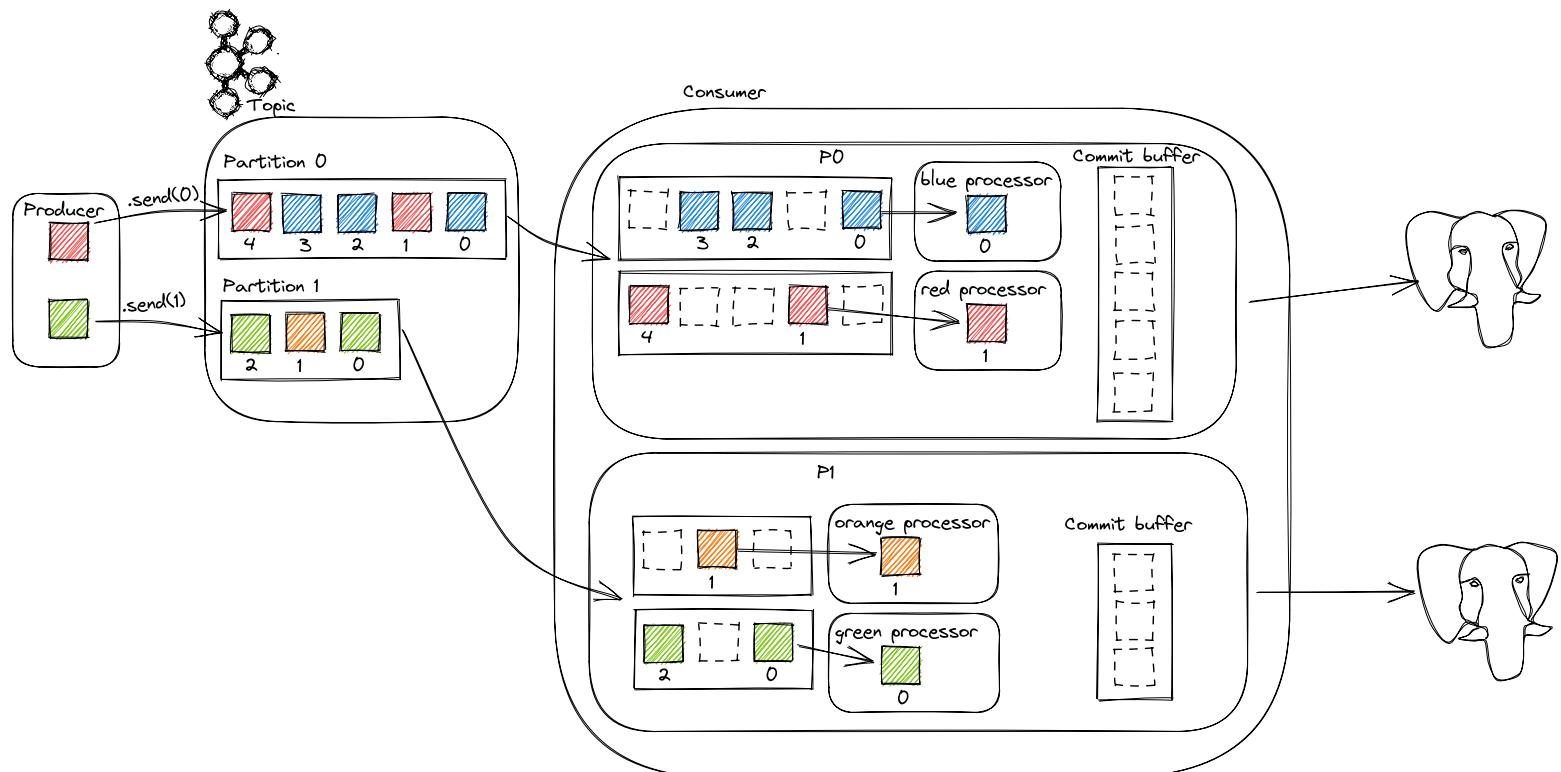
GOALS

- maximize concurrency
- minimize contention
- no concurrent updates to single row

SUBPARTITION QUEUES



ALL TOGETHER



ANOTHER ISSUE

ANOTHER ISSUE

- In-memory queuing

ANOTHER ISSUE

- In-memory queuing
- Memory overloads

ANOTHER ISSUE

- In-memory queuing
- Memory overloads
- Cap on messages in memory

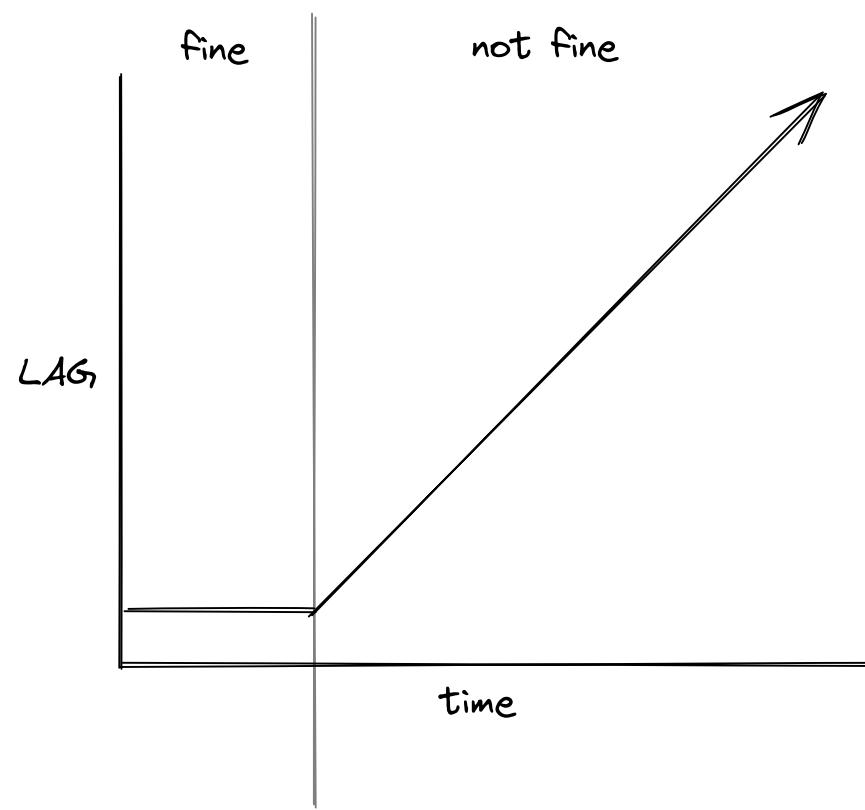
SUDDENLY

everything was fine

SUDDENLY

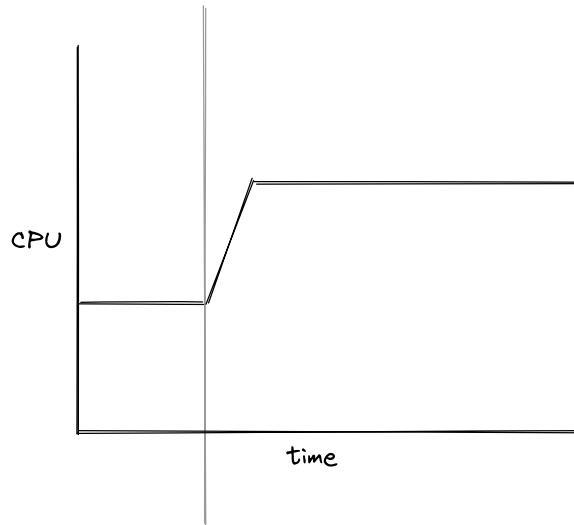
everything was fine

until it wasn't

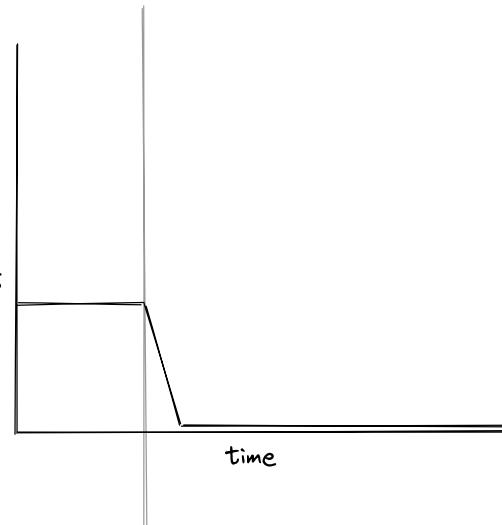


??
?

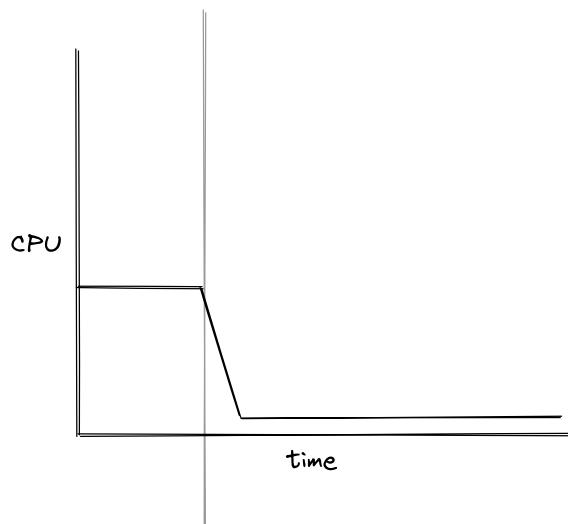
EXPECTATION



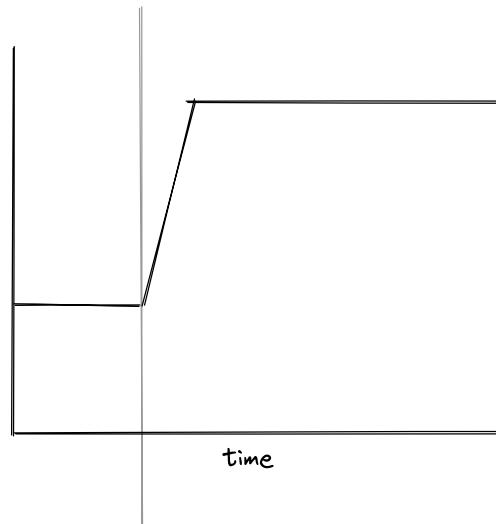
IDLE
CONN



REALITY



IDLE
CONN



OBSERVABILITY

OBSERVABILITY

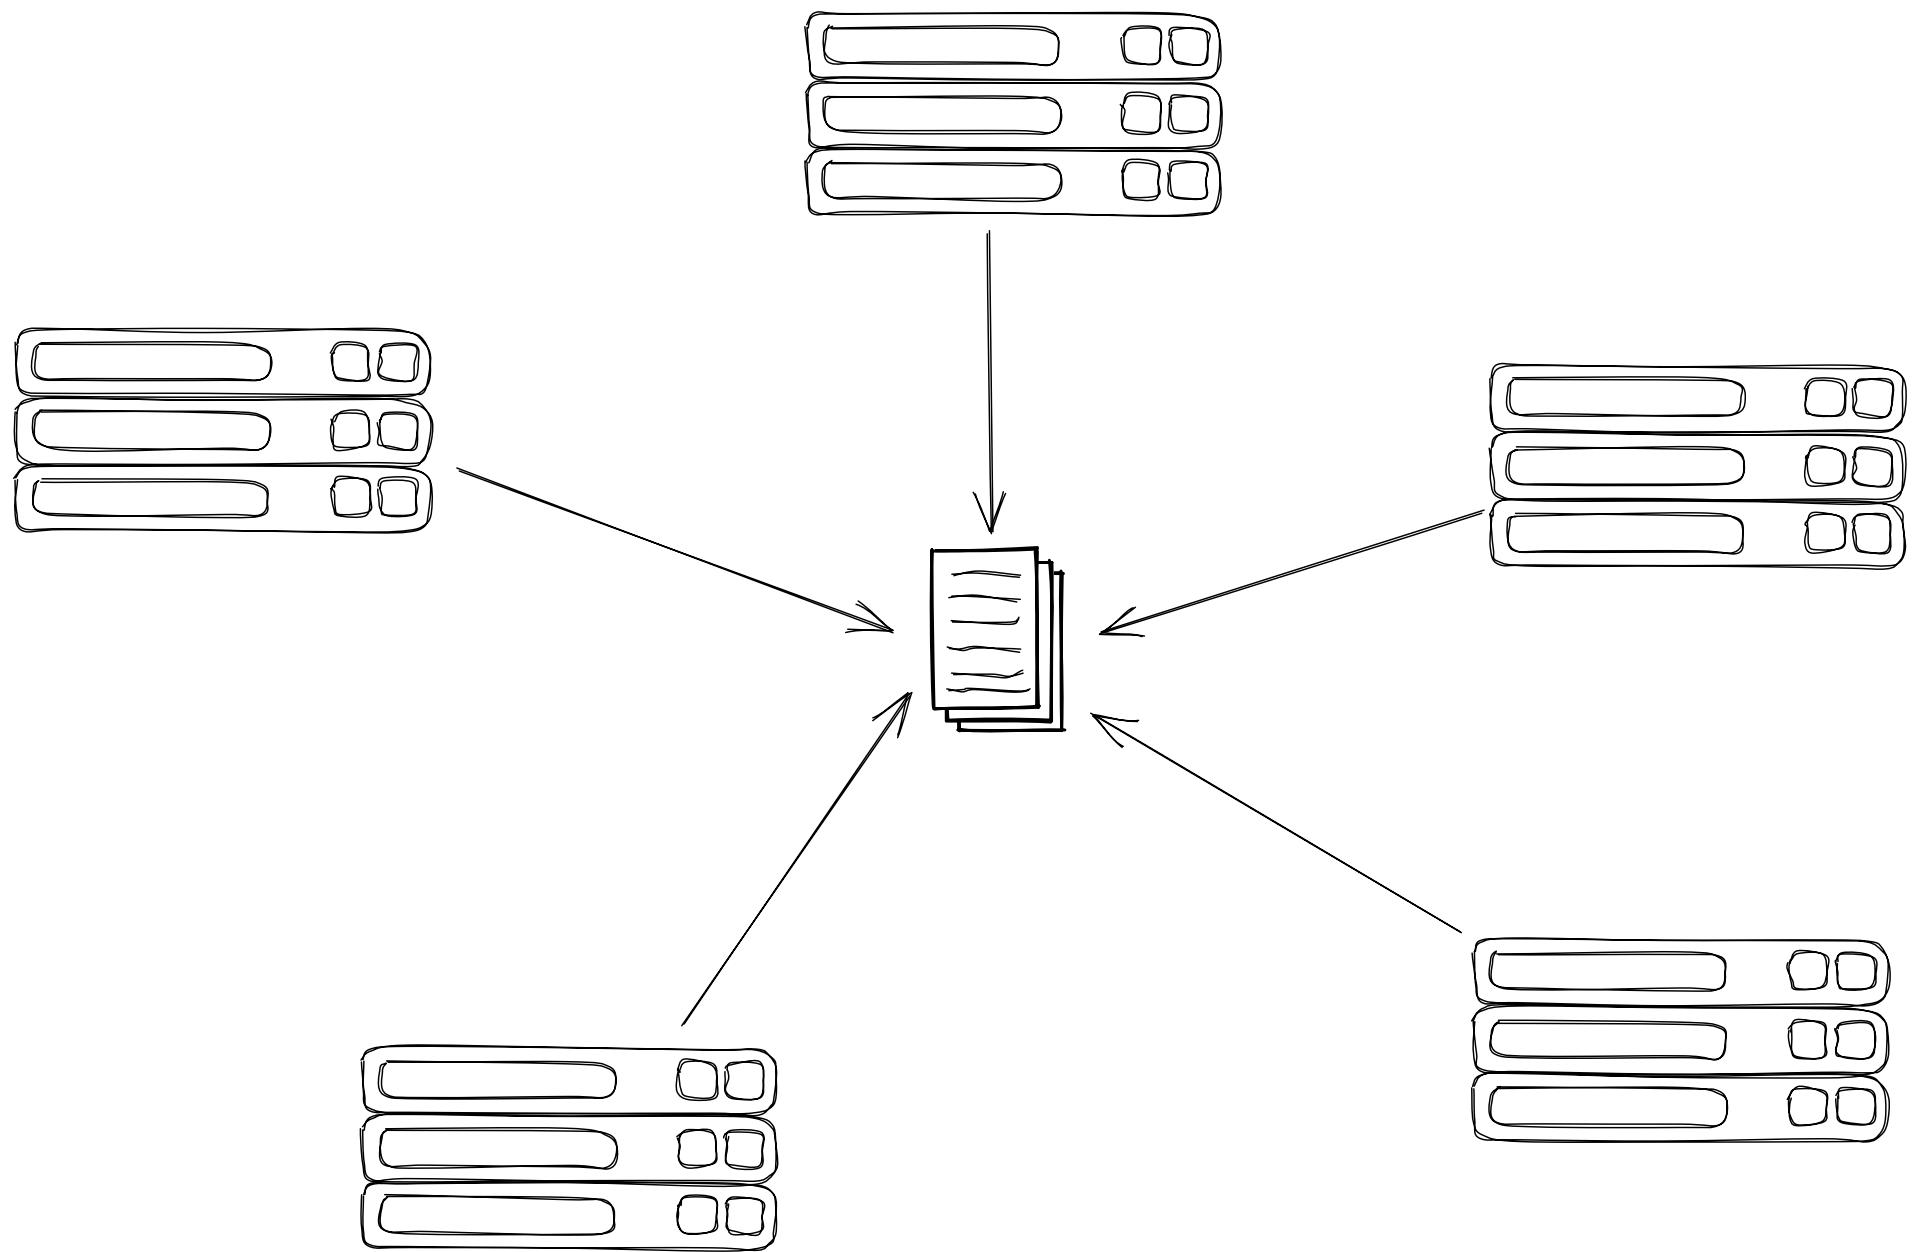
- Only metrics

OBSERVABILITY

- Only metrics
- Unstructured logs on boxes

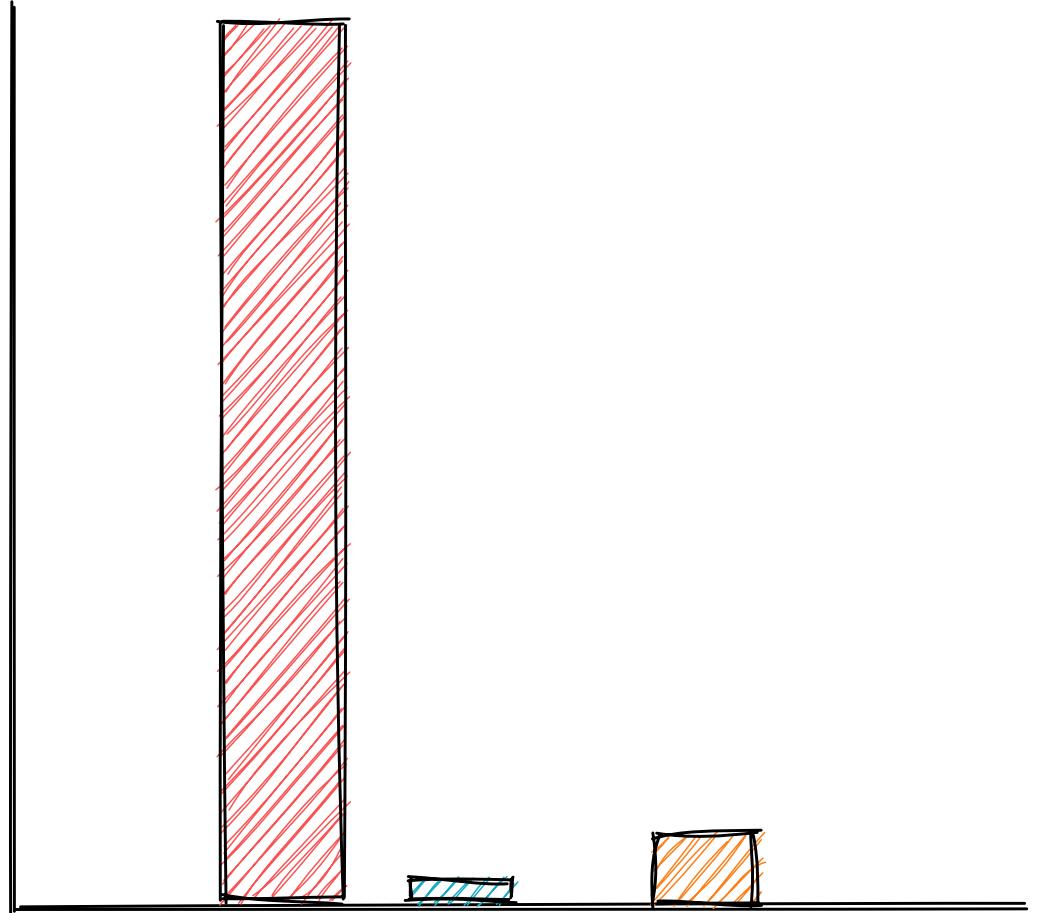
OBSERVABILITY

- Only metrics
- Unstructured logs on boxes
- I insisted on getting centralized logging



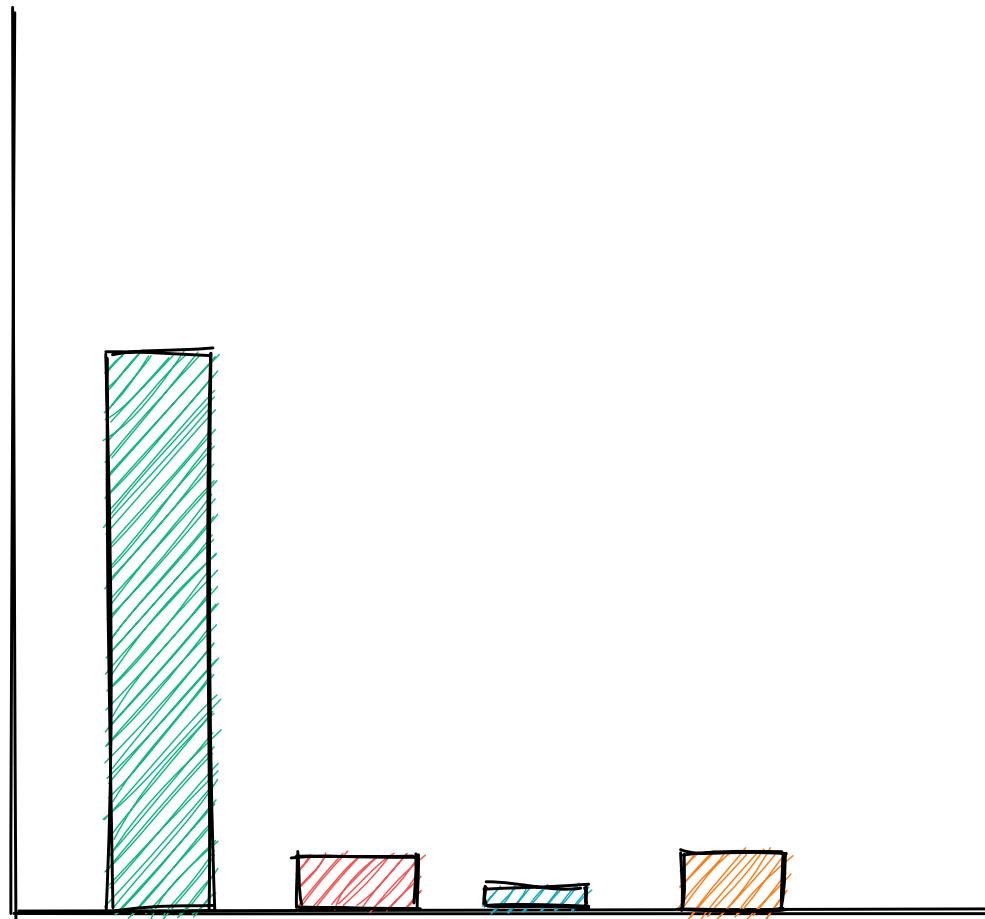
```
{  
    "app_id": "9... ",  
    "subscription_id": "6... ",  
    "sql": "UPDATE ... WHERE id=6... ",  
    "hostname": "consumer-01"  
}
```

Count



clothes.ly
next largest customer
other customers

Count



clothes.ly 51
other clothes.ly
next largest customer
other customers

WHAT?

WHAT?

- Tons of individual updates

WHAT?

- Tons of individual updates
- Incompatible updates

WHAT?

- Tons of individual updates
- Incompatible updates
- Location moving all over



location	color	level	device type	identifier
Chicago	Red	VIP	email	admin@clothes.ly
NYC	Blue	User	email	admin@clothes.ly
Tokyo	Pink	Anon	email	admin@clothes.ly

ONESIGNAL

ONESIGNAL

- More than just push

ONESIGNAL

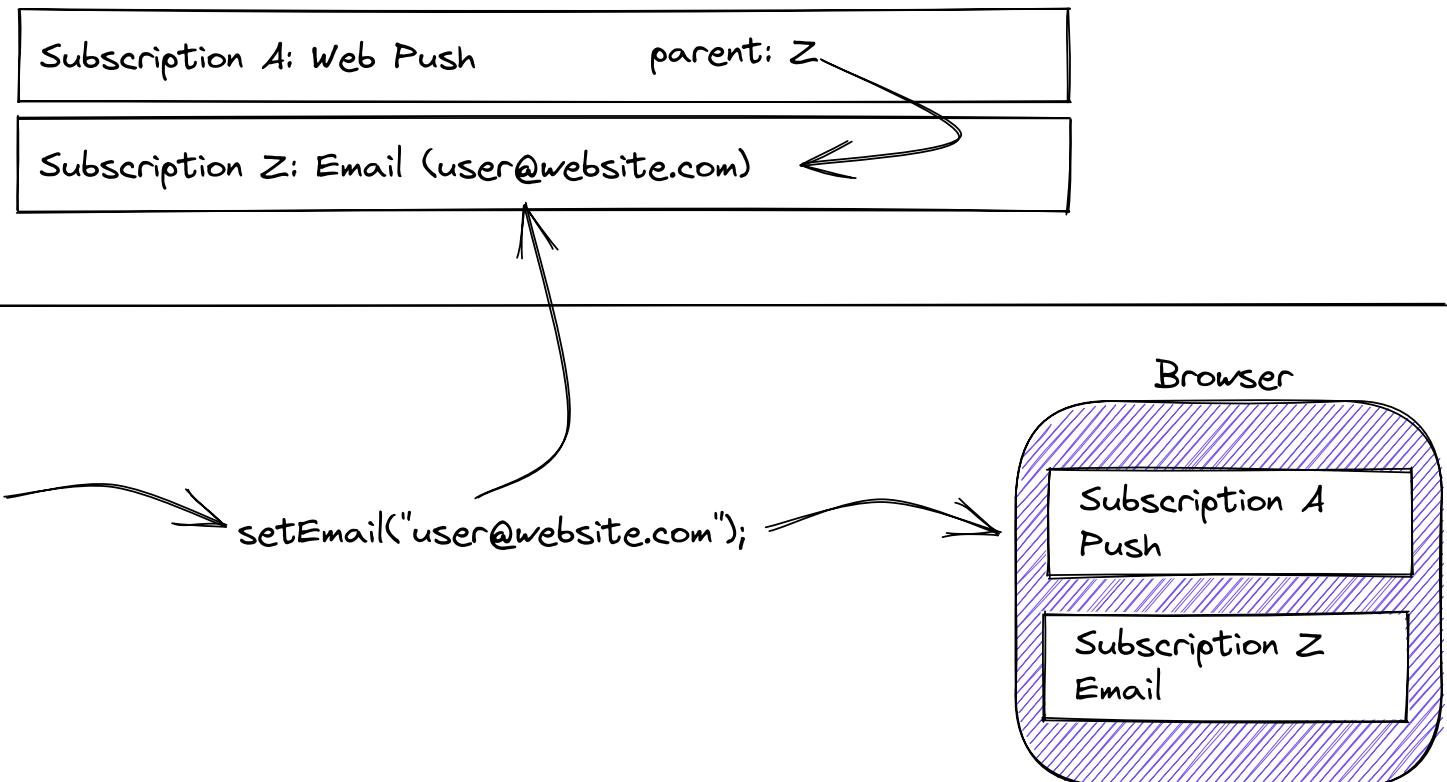
- More than just push
- Omnichannel messaging

ONESIGNAL

- More than just push
- Omnichannel messaging
- Push, email, sms, in-app

setEmail

setEmail



```
COUNT(*) ...
```

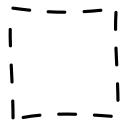
```
5,000,000
```

```
COUNT(*) ... WHERE parent_player_id=S1
```

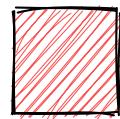
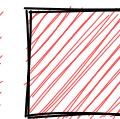
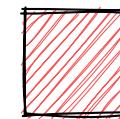
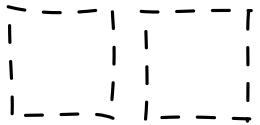
```
4,800,000
```



S1 Update

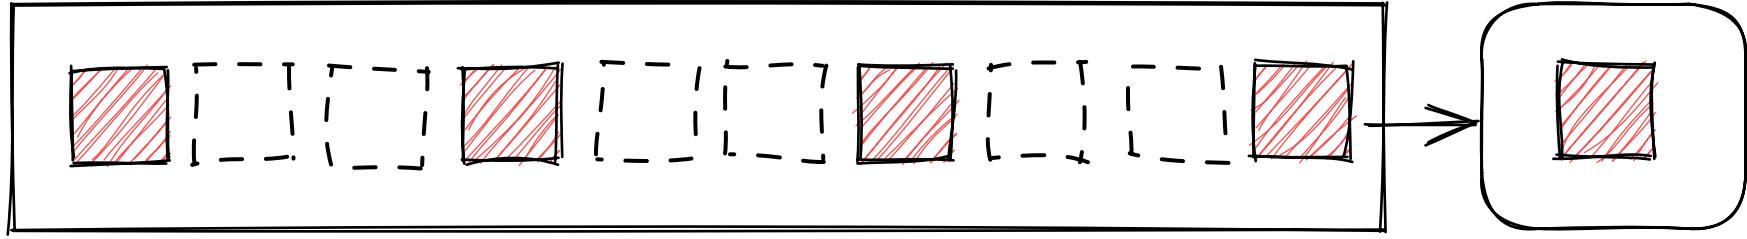


Other Update

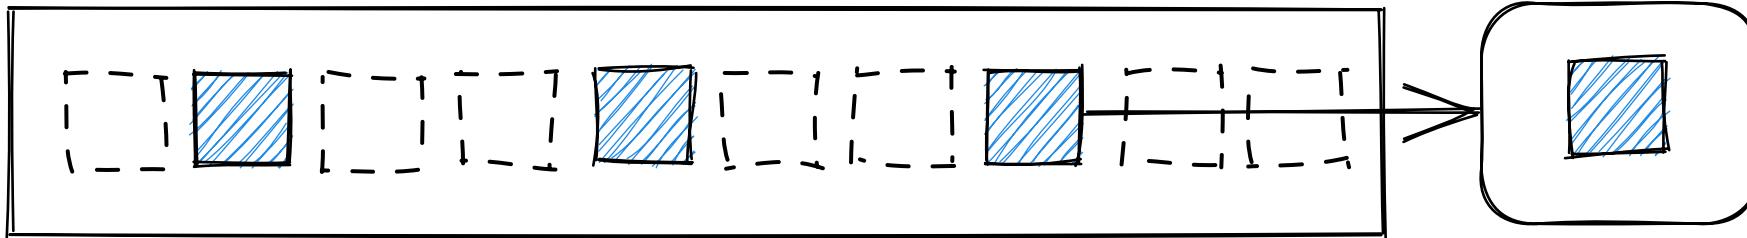


**WHY IS THAT A
PROBLEM?**

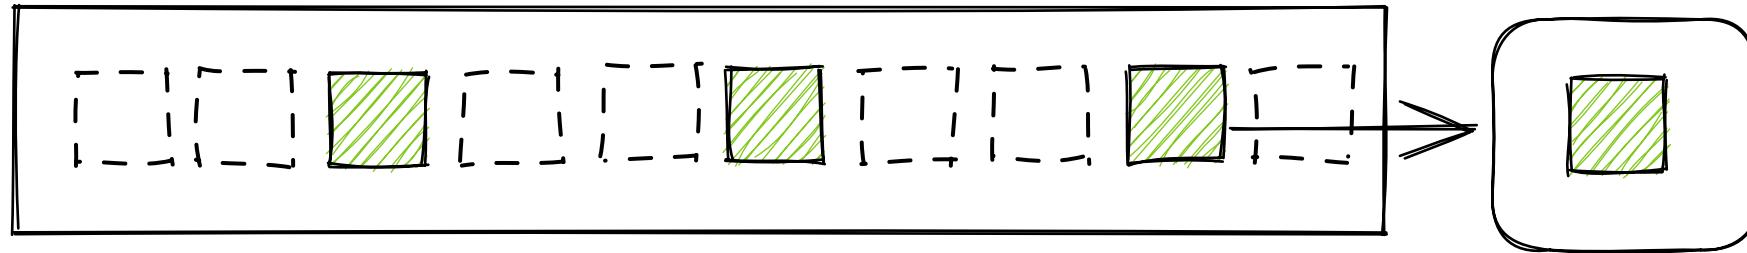
Q_0



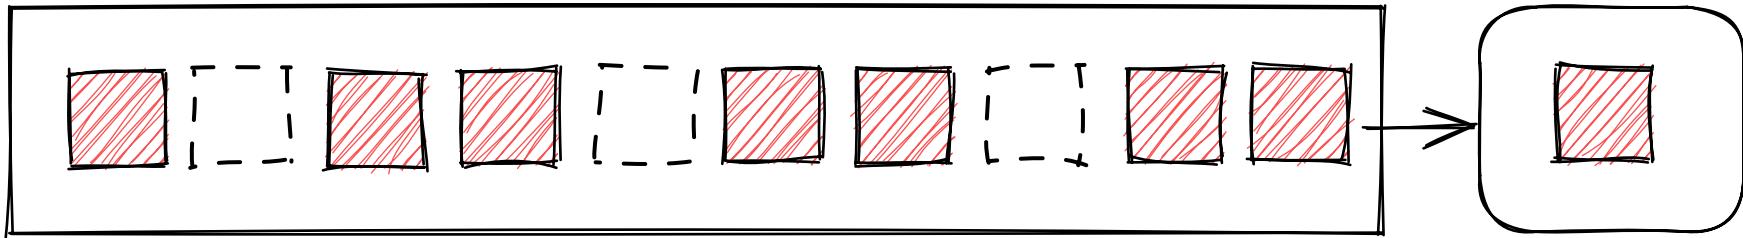
Q_1



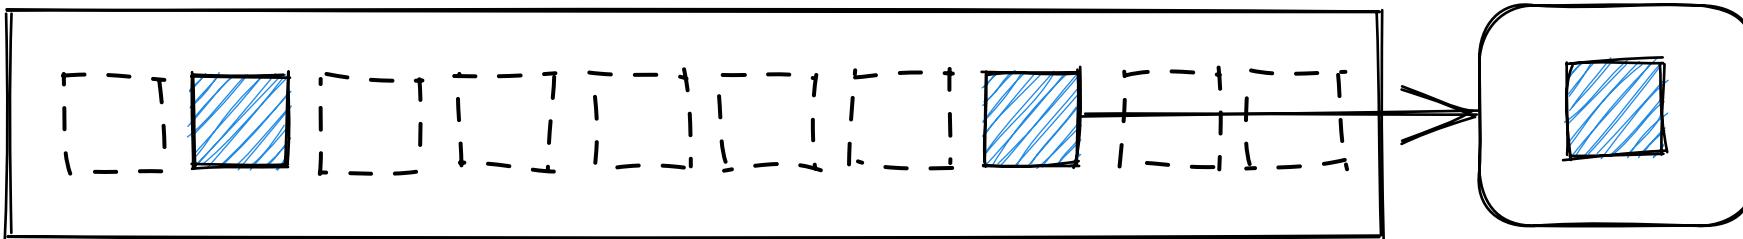
Q_2



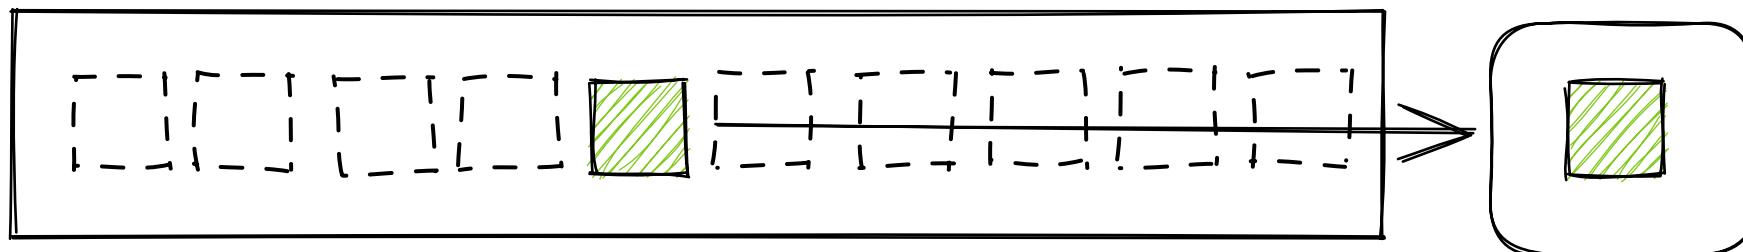
Q_0



Q_1



Q_2



**OK BUT IN REALITY
IT WAS WORSE**

WHAT DID WE DO?

WHAT DID WE DO?

- Skip the updates

WHAT DID WE DO?

- Skip the updates
- Fix message limiting

WHAT DID WE DO?

- Skip the updates
- Fix message limiting
- Limit subscription linking

**WHAT DID WE
LEARN?**

WHAT DID WE LEARN?

- Shift API write workloads to async workers

WHAT DID WE LEARN?

- Shift API write workloads to async workers
- Benefits of subpartition queueing

WHAT DID WE LEARN?

- Shift API write workloads to async workers
- Benefits of subpartition queueing
- Struggles of subpartition queuing

WHAT DID WE LEARN?

- Shift API write workloads to async workers
- Benefits of subpartition queueing
- Struggles of subpartition queuing
- Centralized observability

WHAT DID WE LEARN?

- Shift API write workloads to async workers
- Benefits of subpartition queueing
- Struggles of subpartition queuing
- Centralized observability
- Customers are so creative

LILYMARA.XYZ



